

Pre-Export Row Scaling for Mixed-Integer Programs: Mechanism Attribution and Residual-Budget Contracts

Mathias Rodríguez Castro

Facultad de Ingeniería, Universidad de la República, Montevideo, Uruguay

mathiasr@fing.edu.uy

Abstract

Positive row scaling preserves a mixed-integer program in exact arithmetic while changing its numerical representation. On 306 generated hydrothermal-dispatch models, with one LP basis held fixed, row scaling with geometric-mean local normalization lowers the basis condition number on every model, by six to eight orders of magnitude in the median of each class. Such a gain does not identify which preprocessing information caused it. We show that an apparent block-metadata effect was a coupling-kernel effect: geometric-mean normalization makes an extrema-based block statistic identically one, so a block-relative coupling rule reduces to Euclidean normalization, and a control without ownership information reproduces its spectral gain. After local whitening, an idealized coupling problem has the closed-form optimal factor $(1 + \rho^2)^{-1/2}$, which needs no ownership map. Residual budgets in application units, by contrast, survive consistent row re-emission. We characterize exactly the factors compatible with a budget and coefficient limits, with an incompatibility certificate, the constrained logarithmic minimax factor, the integer-rounding term that scaling cannot remove, and exact binary64 export with power-of-two factors. In reconstructed dispatch runs, scaling used 30–32% less solver work on instances completed by every policy, but one scaled-policy timeout reversed the PAR10 ordering. In an exactly enumerated stress test, adding the budget to a centering rule never lost a verified configuration on any of three solvers, while gains over the unscaled model appeared with HiGHS, not with Gurobi, and not with CPLEX, where tightening the solver tolerances was the effective intervention. Scaling effects are real; their attribution and guarantees must be made explicit.

1 **Keywords:** mixed-integer linear programming; row scaling; numerical reliability; equilibration;
2 residual budgets; computational reproducibility.

3 1 Introduction

4 A model generator can emit the same mixed-integer linear program (MILP) in many mathematically
5 equivalent forms. Multiplying a constraint and its right-hand side (RHS) by a positive number
6 preserves the MILP in exact arithmetic, yet it changes the coefficients on which presolve and LP
7 factorizations operate and the original-unit meaning of an absolute residual tolerance. Solvers scale
8 internally; a generator can also scale before export and use information that a matrix file does not
9 carry, such as row roles, variable ownership and application units [11, 13].

10 The numerical effect of external row scaling can be large. On 306 dispatch models from a
11 hydrothermal generator, holding one LP basis fixed across representations, geometric-mean row
12 scaling lowers the 1-norm condition number of that basis on every model, by six to eight orders of
13 magnitude in the median of each class (Section 7.1). A gain of this kind leaves two questions open.
14 Which information is responsible for it? And what does the chosen representation guarantee once

a solver returns a point? Representation quality, solver effort, failure-sensitive performance and verified accuracy are different endpoints, and evidence for one does not transfer to another.

This paper answers both questions for external row scaling of generated MILPs. Its contributions are the following.

- (A) *Attribution*. A comparison identifies the effect of block ownership only when coverage, row roles, kernels and safeguards are matched. With such controls, the order-of-magnitude conditioning advantage of a block-relative coupling rule over flat geometric-mean normalization is reproduced by a rule that uses no block scales, including its growth with block heterogeneity (Sections 4 and 7.2).
- (B) *Mechanism*. Geometric-mean row centering makes the post-local extrema-based block statistic identically one, so the block-relative coupling formula reduces to Euclidean row normalization. After local whitening, the idealized coupling problem has the closed-form optimal common factor $(1 + \rho^2)^{-1/2}$, which requires no ownership metadata (Section 4).
- (C) *Contract*. Residual budgets stated in application units survive consistent row re-emission. We characterize the admissible row factors exactly, with necessary and sufficient compatibility conditions, an incompatibility certificate, the constrained logarithmic minimax factor, the integer-rounding term that scaling cannot remove, and exact binary64 export with power-of-two factors (Section 5).
- (D) *Validation*. A fixed-basis audit of 306 dispatch models, a fresh audit of 1,080 matched matrix exports, a reconstruction of 240 operational solver runs, and an exactly enumerated stress experiment with 4,752 optimizer calls on three solvers (Sections 6 and 7).

The results progress from the representation to the solver. Scaling changes the numerical representation by orders of magnitude, and the kernel determines how much. The apparent metadata effect is a kernel effect with a simpler explanation. Residual budgets have measurable consequences: adding them to a centering rule never lost a verified stress configuration on any of the three solvers tested. Whether a solver accepts the resulting points, and how much effort it spends, remain properties of the solver pipeline, to be measured and verified in original units.

2 Related work

Diagonal scaling has a long history, including the logarithmic method of Curtis and Reid [4]. Elble and Sahinidis [6] compare scalers for simplex and find that additional complexity need not improve performance. Klotz [16] relates numerical instability to formulation and finite precision. Optimal diagonal preconditioning is a broader problem: Qu et al. [19] use quasiconvex and semidefinite formulations, while Demmel [5] gives block-Jacobi guarantees. Our triangular calculation is a scalar special case of this problem that explains the attribution result.

Robustness is also an established objective in MIP scaling. Berthold and Hendel [2] learn to select between standard and Curtis–Reid scaling in Xpress from an indicator of numerical difficulties. Their implementation uses power-of-two factors and their evaluation examines inconsistencies across seeds. Our external construction instead states row-level admissibility conditions; it does not predict or replace internal solver scaling.

Modeling environments already support generator-level information: Pyomo’s `ScaleModel` applies component factors [18], GAMS has variable and equation scale attributes [7], and IDAES derives and propagates scales across submodels [12, 13]. Plasmio.jl retains local problems and linking constraints [15]. Consequently, using modeling metadata before export is not new. Our identifiability result concerns the loss of one specific block statistic after a preceding normalization.

Tolerance-based scaling also has direct precedent. Gurobi documents absolute feasibility and integrality tolerances and the treatment of very small coefficients [8], and recommends external scaling that makes tolerance-sized violations negligible in application units [11]. Pritchett, Howell and Grebow [17] similarly scale constraints using physical accuracy requirements and SNOPT’s tolerance. Section 5 combines this known transfer with coefficient limits, identifies incompatible specifications, and solves the constrained centering problem. Integer rounding and exact binary export impose additional restrictions.

Integer-aware presolve can instead strengthen the formulation. GCD division and RHS rounding are established reductions [1]; on integer rows they can remove a narrow feasibility margin that pure scaling preserves. We include that distinction in the stress comparison, rather than treating tolerance-based scaling as a replacement for presolve.

Finally, Cheung, Gleixner and Steffy [3] provide certificates for integer-programming results, and Hoen and Gleixner [10] analyze individual branch-and-bound decisions. Our rational residual checks address primal acceptance, and independent enumeration supplies optimality on the small binary instances. What this paper adds to these lines of work is attribution: matched comparisons of actual exports that separate what a scaling policy changes from which information causes the change.

3 Model, policies and comparisons

Consider

$$\min c^T x \quad \text{subject to } Ax \leq b, \quad l \leq x \leq u, \quad x_j \in \mathbb{Z} \ (j \in J),$$

where equalities and opposite inequality senses are allowed. For a positive diagonal matrix D_r , replacing (A, b) by $(D_r A, D_r b)$ preserves every constraint, variable domain and objective value in exact arithmetic. The whole row, including its RHS, is multiplied by one positive number; dividing different coefficients of a row by different block scales is not a row scaling and is outside this equivalence.

The generator provides a partition of the local rows into nonempty blocks, a set of coupling rows, and a column-to-block map β . A scaling policy is then determined by four separate choices: *coverage*, the rows eligible for scaling; the *kernel*, which maps an eligible row to a candidate factor; *safeguards*, which decide whether a candidate is applied; and *information*, the generator data the kernel may read, such as roles, block scales and ownership. An effect can be attributed to information only when the other three choices are matched.

For a nonzero row, put $m_r = \min_{a_{rj} \neq 0} |a_{rj}|$ and $M_r = \max_j |a_{rj}|$. The geometric-mean (GM) factor is $(m_r M_r)^{-1/2}$ and the Euclidean (L2) factor is $1/\|a_r\|_2$. Both are coefficient-only, but they are different kernels, and Table 1 keeps the distinction visible.

Relabeling blocks without changing their assignment is an invariance check, whereas permuting the assignment of unequal block scales tests whether the map is useful on the chosen instances.

4 Identifying the mechanism

This section shows that, under exact geometric-mean local normalization, the extrema-based post-local block statistic used by the original policy carries no block-scale information, so the coupling stage cannot use it, and it derives the simpler mechanism that remains.

Table 1: Policies. Coverage and safeguards are matched across the comparisons in which a policy is used; a role map and a block-scale map are separate information sources. The block scales s_i are defined in [Section 4](#).

Policy	Local rows	Coupling rows	Information used
Base	None	None	None
Flat-GM	GM	GM	Row coefficients
Flat-L2	L2	L2	Row coefficients
Role-Hybrid	GM	L2	Row roles, coefficients
SA-Mat	GM	$\left[\sum_j (a_{rj}/s_{\beta(j)})^2\right]^{-1/2}$	Roles, block scales, ownership
SA-Pre	GM	Same expression, pre-local scales	Roles, original block scales

4.1 Collapse of post-normalization block centers

For each nonzero local row r , write m_r and M_r for its smallest and largest nonzero coefficient magnitudes. The geometric-mean kernel applies $d_r = (m_r M_r)^{-1/2}$. Let $\tilde{m}_r = d_r m_r$ and $\tilde{M}_r = d_r M_r$ denote the exported extrema. For a nonempty set I of local rows, the block statistic is

$$s_I = \sqrt{\left(\min_{r \in I} \tilde{m}_r\right) \left(\max_{r \in I} \tilde{M}_r\right)}.$$

Proposition 4.1 (Exact collapse and the effect of a skip band). *If every row receives its geometric-mean factor, then $s_I = 1$ for every nonempty grouping I of local rows. More generally, suppose the only safeguard replaces d_r by 1 when the candidate lies in $(1/u, u)$, where $u \geq 1$. Then $s_I \in [1/u, u]$, independently of the original coefficient scales and of the grouping of the rows.*

Proof. Under exact normalization, $\tilde{m}_r = \sqrt{m_r/M_r} = 1/\tilde{M}_r$. Consequently $\min_r \tilde{m}_r = 1/\max_r \tilde{M}_r$, proving the first statement. For the second, let $c_r = \sqrt{\tilde{m}_r \tilde{M}_r}$. An accepted row has $c_r = 1$; a skipped row has $c_r = \sqrt{m_r M_r} \in (1/u, u)$. Thus $u^{-2} \leq \tilde{m}_r \tilde{M}_r \leq u^2$ for every row. Choose p attaining the maximum exported row magnitude and q attaining the minimum. Then

$$u^{-2} \leq \tilde{m}_q \tilde{M}_q \leq \tilde{m}_q \tilde{M}_p \leq \tilde{m}_p \tilde{M}_p \leq u^2.$$

Taking square roots gives the result. $\square$

For a coupling row a_r , let $\beta(j)$ assign column j to a local block, and define the block-weighted candidate and its ownership-blind counterpart by

$$\gamma_r^{\text{block}} = \left(\sum_j a_{rj}^2 / s_{\beta(j)}^2\right)^{-1/2}, \quad \gamma_r^{\text{L2}} = \left(\sum_j a_{rj}^2\right)^{-1/2}.$$

Blockless columns may use an arithmetic mean of valid block scales or the unit fallback; both conventions preserve the interval in [Proposition 4.1](#).

Proposition 4.2 (Loss of ownership information). *Under exact local normalization and with no coefficient filtering, $\gamma_r^{\text{block}} = \gamma_r^{\text{L2}}$ for every coupling row, so changing the ownership map cannot change the candidate export; applying the same coupling safeguards to identical candidates preserves equality. Under the skip-band assumptions of [Proposition 4.1](#),*

$$u^{-1} \leq \frac{\gamma_r^{\text{block}}}{\gamma_r^{\text{L2}}} \leq u,$$

and, before coupling safeguards, the condition numbers of the two exports differ by at most a factor u^2 on their common nonzero singular spectrum.

Proof. The equality follows by substituting $s_i = 1$. In the second case,

$$u^{-2}\|a_r\|_2^2 \leq \sum_j a_{rj}^2/s_{\beta(j)}^2 \leq u^2\|a_r\|_2^2,$$

which gives the factor bound. Write the block-based export as EA_{L2} , where E has unit entries on local rows and entries in $[1/u, u]$ on coupling rows. Multiplication by E changes each singular value by a factor between $1/u$ and u and preserves rank. Applying these inequalities to the largest and smallest positive singular values in both directions proves the condition-number bound. $\square$

Both statements assume that the block summary sees every coefficient. Coefficient-window vetoes and filtered extrema lie outside the hypotheses: a row with magnitudes $(10^{-8}, 3, 10^8)$ is geometrically centered, but restricting the summary to $(10^{-6}, 10^6)$ produces $s_I = 3$, a nonunit statistic created by the filter rather than by retained physical scale. The export audit in [Section 7](#) measures how often the implemented safeguards break exact equality.

4.2 What a matched comparison identifies

With equal coverage and equal local kernels, two exports agree on their local rows, so any difference is confined to the coupling rows. This localizes the difference without attributing it to ownership, because the coupling kernels may still differ. The ownership control therefore keeps local GM and coupling L2 normalization but replaces every block scale by one (Role-Hybrid in [Table 1](#)); SA-Mat versus Role-Hybrid is the ownership comparison. Flat-GM changes the coupling kernel and Flat-L2 changes the local kernel, so neither isolates ownership.

Equality also requires equal *applied* factors, not merely an inactive coupling stage in one policy. For local rows $x_1 \leq 1$, $x_2 \leq 1$ and the coupling row $x_1 + 10^{-2}x_2 \leq 1$, both block scales are one and the Euclidean coupling candidate $(1 + 10^{-4})^{-1/2}$ falls inside the skip band $(1/2, 2)$, whereas the flat GM candidate 10 is applied. Both policies cover every row, yet their exports differ.

4.3 The spectral slope in a stylized model

Take $B \geq 2$ local singleton rows with positive coefficients $v_i = 10^{e_i}$, where $\min_i e_i = -R$ and $\max_i e_i = R$, and one coupling row $v^\top$. Exact local normalization produces I_B and unit post-local block centers. A whole-row coupling factor γ gives

$$A_\gamma = \begin{bmatrix} I_B \\ \gamma v^\top \end{bmatrix}, \quad A_\gamma^\top A_\gamma = I_B + \gamma^2 v v^\top, \quad \kappa_2(A_\gamma) = \sqrt{1 + \gamma^2 \|v\|_2^2}.$$

Flat geometric-mean normalization selects $\gamma = 1$, whereas both the post-local block rule and the hybrid Euclidean control select $\gamma = 1/\|v\|_2$. Hence

$$\kappa_2(A_{\text{GM}}) = \sqrt{1 + \sum_i 10^{2e_i}}, \quad \kappa_2(A_{\text{block}}) = \kappa_2(A_{\text{hybrid}}) = \sqrt{2}.$$

For fixed B the ratio is of order 10^{-R} : the advantage grows with block heterogeneity and survives removal of the ownership map. Every coupling coefficient is multiplied by the same γ , so all within-row ratios are preserved; dividing each v_i by its own 10^{e_i} would change the constraint and is not a row scaling.

4.4 Optimal coupling factor after local whitening

The stylized model suggests a constructive replacement for the metadata explanation: once local scale has been removed, choose the coupling factor from the coupling coefficients themselves. In the idealized triangular model this choice has an exact solution.

Proposition 4.3 (Optimal common coupling factor after local whitening). *Let $R \in \mathbb{R}^{m \times n}$, with $m, n \geq 1$, and set $\rho = \|R\|_2$. For*

$$C_\gamma = \begin{bmatrix} I_n & 0 \\ \gamma R & \gamma I_m \end{bmatrix}, \quad \gamma > 0,$$

the unique factor minimizing $\kappa_2(C_\gamma)$ is

$$\gamma_* = (1 + \rho^2)^{-1/2}, \quad \min_{\gamma > 0} \kappa_2(C_\gamma) = \sqrt{1 + \rho^2} + \rho.$$

On an admissible interval $0 < \ell \leq \gamma \leq h$, the optimum is $\min\{h, \max\{\ell, \gamma_\}\}$.*

The proof (Online Resource 1) reduces the spectrum by an SVD to two-by-two blocks. For the block at ρ , the singular-value ratio k satisfies $k + k^{-1} = \gamma^{-1} + \gamma(1 + \rho^2)$, whose unique minimizer gives the result.

For a single coupling row, γ_* is the Euclidean normalization of the complete row, including the coefficient of its aggregate variable. The idealized coupling problem therefore has a simple optimal whole-row normalization that uses no ownership information, and this is the mechanism the matched control isolates empirically. The statement concerns a common factor on a locally whitened triangular structure; for general generated MIPs it motivates the Euclidean coupling kernel rather than certifying it.

Why conditioning needs a residual bound. The interval form of Proposition 4.3 matters. In A_γ , κ_2 decreases to one as $\gamma \rightarrow 0$, while a fixed residual on the exported coupling row corresponds to an original residual larger by the factor $1/\gamma$. A conditioning criterion alone cannot say how small a coupling factor may be: on any interval its minimizer for A_γ is the lower endpoint ℓ . Section 5 derives that endpoint from an application-unit residual budget.

5 Residual-budget contracts

Block ownership is one kind of generator information; a tolerance stated in application units is another. A residual budget says how much violation of an original constraint is acceptable. It must be supplied by the application or declared as a numerical test parameter; it cannot be inferred from a block label or a zero right-hand side. Scaling rows so that solver tolerances correspond to negligible application-unit violations is established practice [8, 11, 13]. The results below make its combination with coefficient limits exact: they characterize every admissible factor, certify when none exists, select the centered one, and identify what row scaling cannot control.

For a nonconstant row $a^T x \leq b$, define the original residual $v(x) = (a^T x - b)_+$ and a budget $\delta > 0$. Let $\epsilon > 0$ bound the residual of the *exported* row. This is a hypothesis about the returned point rather than a consequence of a solver parameter: presolve, internal scaling and finite precision intervene between the two. For a positive row factor α ,

$$(\alpha a^T x - \alpha b)_+ = \alpha v(x) \tag{1}$$

in exact arithmetic, so an exported residual at most ϵ implies $v(x) \leq \delta$ whenever $\alpha \geq \epsilon/\delta$. Absolute residuals give the same identity for equalities. Write $m = \min_{a_j \neq 0} |a_j|$ and $M = \max_j |a_j|$, and let $0 < \ell \leq U$ limit the exported coefficient magnitudes. An additional RHS limit U_b intersects the upper endpoint below with $U_b/|b|$ for $b \neq 0$.

Proposition 5.1 (Admissible factors and constrained centering). *The factors that place every nonzero coefficient in $[\ell, U]$ and satisfy $\epsilon/\alpha \leq \delta$ form exactly the interval*

$$\mathcal{I} = [L, H], \quad L = \max\{\epsilon/\delta, \ell/m\}, \quad H = U/M. \quad (2)$$

The interval is nonempty if and only if

$$M/m \leq U/\ell \quad \text{and} \quad \delta \geq \epsilon M/U. \quad (3)$$

When it is nonempty, the unique minimizer of

$$\min_{\alpha \in \mathcal{I}} \max_{a_j \neq 0} |\log(\alpha |a_j|)| \quad (4)$$

is

$$\alpha^* = \min\{H, \max\{L, (mM)^{-1/2}\}\}. \quad (5)$$

Proof. The coefficient bounds are equivalent to $\alpha m \geq \ell$ and $\alpha M \leq U$. Together with $\alpha \delta \geq \epsilon$ they give (2). The condition $L \leq H$ is equivalent to the two inequalities in (3). To solve the objective, put $t = \log \alpha$, $q = (\log m + \log M)/2$ and $h = (\log M - \log m)/2$. Since the extreme absolute logarithms occur at an endpoint,

$$\max_j |t + \log |a_j|| = \max\{|t + \log m|, |t + \log M|\} = h + |t + q|.$$

The unique minimizer on $[\log L, \log H]$ is the projection of $-q$ onto that interval. Exponentiating gives (5). $\square$

Incompatibility certificate. When $L > H$, (3) names the violated requirement. A within-row spread $M/m > U/\ell$ cannot be repaired by any row factor; a budget $\delta < \epsilon M/U$ conflicts with the coefficient ceiling even for a perfectly balanced row. In both cases the certificate concerns the numerical specification, not the optimization model: exporting the unscaled row does not meet the contract, and resolving the conflict requires revising units, limits or the formulation. The factor α^* optimizes coefficient centering subject to the contract; it is not a condition-number minimizer.

Corollary 5.2 (Invariance under consistent re-emission). *Replace (a, b, δ) by $(qa, qb, q\delta)$ for any $q > 0$, keeping ϵ, ℓ, U fixed. The interval becomes $\mathcal{I}/q$, abstention is unchanged, and the selected factor becomes α^*/q . The exported row and the budget-normalized original residual $v(x)/\delta$ are therefore unchanged.*

Proof. Both coefficient extremes and the original budget acquire the factor q . Each term defining L , H and $(mM)^{-1/2}$ is divided by q . Equation (1) gives the residual statement. $\square$

A budget is therefore information that survives preprocessing: generators that emit the same constraint in different units, each with the budget expressed in its own units, obtain the same exported row and the same contract. The invariance is lost if a row is re-emitted without transforming its budget, and a fixed cutoff on the factor itself does not have it. Restricting factors to powers of two [2] retains it for power-of-two re-emissions, subject to the representability conditions of Proposition 5.3.

5.1 The integer-rounding term

Let J be the integer columns and let $\bar{x}$ be obtained by rounding only these coordinates of x . Suppose $|\bar{x}_j - x_j| \leq \eta_j$ for $j \in J$ and put $E = \sum_{j \in J} |a_j| \eta_j$. The triangle inequality and (1) give

$$v(\bar{x}) \leq \epsilon/\alpha + E, \quad (6)$$

also for equalities with absolute residuals. Changing the row factor reduces the exported-residual term but leaves the rounding term, which is measured in original units, unchanged. Normalizing a big- M row therefore cannot remove the effect of an integrality tolerance on that row.

For a contract on the rounded assignment, Proposition 5.1 applies with δ replaced by $\delta - E$ when $\delta > E$; if $\delta \leq E$, no positive factor certifies it. The bound is a worst case: actual rounding errors may be smaller or cancel, and other rows can exclude the worst case, so an empty interval shows that the stated information is insufficient rather than that returned points will violate the budget. Under re-emission E becomes qE , so invariance is retained. A tighter integrality tolerance or a reformulation can reduce E ; row scaling cannot.

5.2 Exact export on a dyadic grid

A factor in the real admissible interval does not guarantee an exactly proportional exported row: the products $\text{fl}(\alpha a_j)$ and $\text{fl}(\alpha b)$ are rounded independently, and rounding a computed endpoint to nearest can leave the closed interval. Our arbitrary-factor implementation forms the endpoints as exact rational numbers in the supplied binary64 inputs, rounds them *inward*, and abstains if no positive binary64 factor remains. This certifies admissibility of the factor, not exactness of the products.

Powers of two avoid multiplication rounding. The following intersection also accounts for subnormal numbers and for the RHS. Assume finite binary64 input data and IEEE arithmetic with gradual underflow [14]. Write each nonzero datum $z \in \{a_1, \dots, a_n, b\}$ uniquely as $z = \text{sign}(z) N_z 2^{p_z}$ with N_z a positive odd integer, and put $q_z = \lfloor \log_2 |z| \rfloor$. Let $[L, H]$ be the desired real factor interval, including any rounding allowance from (6).

Proposition 5.3 (Exact dyadic admissibility). *A factor $\alpha = 2^k$ is positive and representable in binary64, belongs to $[L, H]$, and multiplies every nonzero coefficient and the RHS exactly into a finite binary64 value if and only if the integer k satisfies*

$$K_- = \max \left\{ -1074, \lceil \log_2 L \rceil, \max_z (-1074 - p_z) \right\} \leq k, \\ k \leq K_+ = \min \left\{ 1023, \lfloor \log_2 H \rfloor, \min_z (1023 - q_z) \right\}.$$

If $K_- > K_+$ there is no such factor. Otherwise an exponent nearest to $-\frac{1}{2} \log_2(mM)$ among the integers in $[K_-, K_+]$ minimizes the logarithmic coefficient objective (4) on this exact-export set; a tie may be resolved by choosing the larger factor.

Proof. The factor itself requires $-1074 \leq k \leq 1023$. Since the datum's odd significand is unchanged, the product $N_z 2^{p_z+k}$ lies on the subnormal grid exactly when $p_z + k \geq -1074$; its highest bit stays below overflow exactly when $q_z + k \leq 1023$. The original significand has at most 53 bits, so these two conditions are also sufficient. Intersecting them over the row data and with $L \leq 2^k \leq H$ gives the interval of integer exponents. The objective is $\frac{1}{2} \log(M/m) + |k \log 2 + \frac{1}{2} \log(mM)|$, proving the nearest-exponent claim. $\square$

Table 2: Endpoints and how they are measured.

Endpoint	Measurement	Results
Exact equivalence	Whole-row proportionality of named exports, including RHS	§7.1, §7.2
Numerical representation	κ_1 of one fixed basis; κ_2^+ of exported matrices	§7.1, §7.2
Mechanism attribution	Matched exports: byte identity and paired spectral ratios	§7.2
Solver effort	Gurobi work units on runs completed by every policy	§7.3
Failure-sensitive performance	PAR10 over a fixed instance pool	§7.3
Solver acceptance	Solver status and returned point	§7.4
Verified correctness	Exact original residuals, bounds and integrality; enumerated optimum	§7.4

The implementation obtains the endpoint exponents from integer ratios and bit lengths, compares candidate objectives exactly as rational numbers, and checks the rounded products against exact rational products. Writing the row to a text file additionally requires a round-trip decimal serializer, since fixed decimal places can discard small values even when the multiplication is exact.

Verification after solving. The contract is conditional on the exported residual, so returned points are checked against the original rows, bounds and integrality restrictions. On the stress instances of Section 6 the check is exact: row residuals are rational expressions in the stored binary64 data, and optima are obtained by enumeration. On the dispatch models the archived checker uses floating-point residuals, which are diagnostics; certifying branch-and-bound results is a stronger task [3].

6 Computational design

The experiments measure the endpoints in Table 2. Each result is stated for its own endpoint and none is used as a proxy for another: a condition number describes a representation, not solver effort or the correctness of a returned point, and a solver status is not an independently verified solution.

6.1 Fixed-basis conditioning audit

A condition number read along each policy’s own solve compares different vertices as well as different representations. This audit holds the vertex fixed. For each of the 306 archived dispatch inputs (204 Simple, 68 SG-Ter-Mer and 34 Full models), the C++ generator exports Base, Flat-GM, Flat-L2, Role-Hybrid and SA-Mat without calling an optimizer. HiGHS 1.15.1 solves the Base LP relaxation with the simplex method on one thread, and its optimal basis is recorded by row and column names. The same basis is then assembled from every policy’s own export: structural basic columns come from the scaled matrix and basic slacks enter as unit columns, as a solver holds them. Before assembly, each export is checked to be a whole-row positive scaling of Base, including RHS values, with identical bounds, objective and integrality (relative tolerance 2×10^{-13}).

The diagnostic is $\kappa_1(B) = \|B\|_1 \|B^{-1}\|_1$, with $\|B^{-1}\|_1$ estimated from a sparse LU factorization by the block 1-norm estimator [9]. Because κ_1 is not invariant under row scaling, the audit measures the conditioning of the exported representation before any solver-internal scaling, not an intrinsic property of the vertex. An independent script rebuilds every basis from the LP text and recomputes each estimate through an exact power-of-two row equilibration, and a precision check bounds κ_1 from both sides on a stress sample. The instance set is complete, and all exports, bases and hashes are archived.

6.2 Matrix attribution audit

A C++ generator of block-structured instances builds local blocks and coupling rows and retains row roles and column ownership, so every policy in Table 1 is available in one implementation and exports can be compared before any optimizer is called. The archived headline cell has four heterogeneous blocks, coefficient exponent range $[-2, 2]$, coupling disturbance severity 6 and 30 seeds. A fresh replication crosses ten seeds (20260909–20260918) with exponent ranges $[-R, R]$ for $R \in \{1, 2, 3\}$, two severities and three coupling patterns: 180 instances and 1,080 exports, written with a serializer that preserves binary64 values on readback. Identity claims rest on direct comparisons of named exports. The spectral pseudo-condition number $\kappa_2^+(A) = \sigma_{\max}(A)/\sigma_{\min}^+(A)$ uses a numerical rank fixed per instance from the row-L2-normalized Base (relative threshold 10^{-12}) and shared by all its variants.

6.3 Operational reconstruction

The operational campaign runs Base, Flat-GM, Flat-L2 and Role-Hybrid on 15 fixed instances from each of the Simple and SG-Ter-Mer classes with Gurobi at relative MIP gaps of 1% and 0.1% (seed 42, 16 threads, 1,800-second limit): 240 assigned runs, whose native logs and input hashes were retrieved from the cluster. The subsets were fixed before outcomes were inspected. Nine of the 15 SG-Ter-Mer instances, including instance077, carry a scenario defect that makes them unfaithful market cases without affecting their validity as MILPs or the comparison of representations of one model (Online Resource 1). The pre-specified primary endpoint is failure-sensitive performance: the arithmetic-mean PAR10 over all 15 instances of a cell, with 18,000 seconds assigned to a non-completion. Solver effort, the secondary endpoint, is Gurobi’s work measure, recovered from the native logs at a common resolution of 0.01 units, on the instances that every policy completed; paired geometric-mean ratios carry exploratory bootstrap intervals (Online Resource 1). The two gap settings reuse the same instances and are not independent replications.

6.4 Solver-robustness campaign

The reconstruction above fixes one solver, one seed, one thread count and 15 instances per class. A second campaign separates what survives changing those choices. Its protocol was frozen before any evaluation run: the design document, the instance split, the run tables and their hashes are stored in the repository, and every planned run is reconciled with its record in a closing audit (Online Resource 1). Reruns are admissible only on documented external failure — node loss, corrupted file, unavailable licence, scheduler kill unrelated to the solver. A solver timeout is an outcome, not a failed job.

Instances. Three families of 500 instances each — Simple, SG-Ter-Mer and Full — form the benchmark. Instances are ordered by the digest of the campaign identifier, the family and the instance file; the first 10, 10 and 30 form a calibration set used only to choose time limits and excluded from every comparison, and the next 490, 490 and 120 form the evaluation set. The Full family is smaller because its runs are the most expensive and the Gurobi licence sustains two concurrent sessions; the reduction was decided from the historical cost model before any policy ran on the pool. A separate pool of 10,000 Simple instances exists and is not used: enlarging the easiest family would inflate the run count without addressing the open question.

Time limits. The family time limits follow a rule registered before the calibration runs: the smallest limit in $\{1800, 3600, 7200\}$ seconds for which the unscaled policy completes at least 80% of

calibration instances at a 1% gap and at least 70% at 0.1%, for Gurobi and for CPLEX, and 7,200 seconds if none qualifies. A non-completion is scored at ten times its own family limit, and PAR10 is never averaged across families.

Solvers. Gurobi 13.0.1, CPLEX 22.1.2 and HiGHS 1.15.1 solve the same exported models through their C++ interfaces, each on one thread, with a primal feasibility tolerance of 10^{-6} , an integrality tolerance of 10^{-5} , and presolve, cuts and internal scaling at their defaults. The thread count was fixed before any comparison; conditioning instrumentation is off so that it cannot perturb the measured effort. Effort measures are solver-specific — work units, deterministic ticks, nodes — and are compared only within a solver. A second registered rule restricts HiGHS to the replication subset in any family where it completes fewer than half of the calibration instances at the family limit; its censoring is reported as a result rather than hidden by a larger limit.

Endpoints and replication. The primary endpoints are PAR10 over the assigned pool and the completion, timeout and censored-gap rates, reported for each family, solver and gap separately; the 1% and 0.1% settings are never pooled. The secondary endpoint is the paired ratio on the instances that every policy completes. Every returned point is verified against the original model in its own units. A replication subset of 50, 50 and 20 instances, fixed before any result was seen, is rerun with four further solver seeds; its inference resamples instances, not instance–seed pairs, because seeds of one instance are not independent replications. Instance077, which reversed the PAR10 ordering in the reconstruction, stays in the evaluation set and is additionally audited on its own, in both its historical and its corrected version, under the historical thread count as well as the campaign’s.

6.5 Enumerated residual-budget stress experiment

Three binary families with 12 variables contain a cardinality row, three positive packing rows and three rows with signed integer coefficients; profits are positive integers. With integer matrix W and integer capacities k , the represented constraints are

$$10^e Wx \leq 10^e k - 0.1 \mathbf{1}, \quad x \in \{0, 1\}^{12}, \quad e \in \{0, 3, 6, 9\}. \quad (7)$$

For these exponents the binary feasible set is exactly $Wx \leq k - \mathbf{1}$, and the stored binary64 right-hand sides are checked to separate the adjacent scaled integer activities. The exponent therefore changes the representation and the continuous relaxation without changing the binary optimum. Because the absolute margin 0.1 is fixed, this is not pure re-emission: it deliberately places infeasible integer points at a controlled distance from the continuous boundary. The budget is 0.01 in the units of the represented row, a declared test parameter.

Eight seeds (100–107), reserved after a one-seed pilot, give 96 models. Each is solved with presolve on and off by HiGHS 1.15.1, Gurobi 13.0.2 and CPLEX 22.1.2 on one thread, with zero target gap, a ten-second limit, and primal and integrality tolerances 10^{-6} . Base, GM and L2 scale every row by their kernels without an identity band; GM-tight uses GM with both tolerances tightened to 10^{-9} ; Budget-only uses the direct factor ϵ/δ_r with $\epsilon = 10^{-6}$; Budget-GM applies [Proposition 5.1](#) with coefficient limits $[10^{-9}, 10^9]$; Budget-round additionally reserves $E_r = 10^{-6} \sum_j |a_{rj}|$ and abstains on the whole model if any row’s interval is empty. Dyadic-GM and Budget-dyadic use power-of-two factors under [Proposition 5.3](#), without and with the budget. These two arms were added after an endpoint rounding issue was found in the initial run and are exploratory; the original seven arms were rerun with the corrected implementation and unchanged seeds. The nine policies give 1,728 planned configurations per solver; Budget-round abstains on 144 of them, leaving 1,584 calls per solver and 4,752 in total.

Table 3: Fixed-basis conditioning audit on all 306 dispatch models. One optimal basis of the Base LP relaxation is assembled from every policy’s own export. The first row is the median κ_1 of the Base basis; the others give the median per-instance ratio $\kappa_1(\text{policy})/\kappa_1(\text{Base})$ and, in parentheses, the number of instances with a ratio below one. Ranges are in Online Resource 1, Table S1.

Policy	Simple ($n = 204$)	SG-Ter-Mer ($n = 68$)	Full ($n = 34$)
Base, median κ_1	8.8×10^{18}	2.9×10^{20}	4.3×10^{20}
Flat-GM	2.4×10^{-7} (204)	7.3×10^{-9} (68)	3.5×10^{-7} (34)
Flat-L2	1.57 (32)	1.51 (10)	1.49 (2)
Role-Hybrid	2.4×10^{-7} (204)	7.3×10^{-9} (68)	3.5×10^{-7} (34)
SA-Mat	2.4×10^{-7} (204)	7.3×10^{-9} (68)	3.5×10^{-7} (34)

A configuration is *verified* when the returned point is within 10^{-9} of a binary assignment that is exactly feasible for the stored binary64 data and attains the optimum found by enumerating all 2^{12} assignments, and the bounds hold within 10^{-9} . Residual budgets and integrality errors are recorded separately, and the rule does not use the solver status. The instances are small and deliberately close to integer feasibility boundaries, so counts describe a failure mechanism rather than practical failure rates, and repeated solver settings on one model are not treated as independent samples. A later control on the same 96 stored models applies GCD division with RHS flooring, an established integer-aware reduction that preserves every binary feasible set but changes the LP relaxation; Online Resource 1 reports it together with two GM variants.

7 Results

7.1 Scaling materially changes numerical representation

All 1,530 exports pass the whole-row scaling check, with largest relative deviation 2.5×10^{-16} , and every policy is evaluated on the same basis. The three policies with GM local normalization lower κ_1 on every one of the 306 models, with median ratios of 2.4×10^{-7} in Simple, 7.3×10^{-9} in SG-Ter-Mer and 3.5×10^{-7} in Full (Table 3). Because the vertex is common, these are differences between representations rather than between the vertices that different solves reach. The independent reconstruction reproduces all 1,530 bases exactly. Because the unscaled values exceed the reciprocal of machine precision, their accuracy was also checked directly on a stress sample of 24 bases from eight models: in each class, the model with the largest Base estimate, the model with the weakest Flat-GM ratio and one seeded random model, under Base, Flat-GM and Flat-L2. A lower bound certified with exact rational residuals and an a posteriori upper bound over all columns of B^{-1} enclose every κ_1 within a relative width of 6×10^{-3} , and the resulting upper bounds on the Flat-GM/Base ratio agree with the estimated ratios to 2.1×10^{-5} (Online Resource 1, Table S2). The median separation of six to eight orders of magnitude, present on every model, is therefore a property of the representations and not of floating-point error in the unscaled bases. An archived computation on a different reference basis, the LP with integers fixed at a Gurobi incumbent, points in the same direction with smaller magnitudes, but its exports were not retained and it is not used as evidence (Online Resource 1).

The kernel, with the safeguards acting on its candidates, determines this separation. Flat-L2 covers the same rows as Flat-GM, yet its median ratio stays between 1.49 and 1.57 and it lowers κ_1 on only 44 of the 306 models. Every row of these models has a coefficient of magnitude at least one, so the Euclidean factor $1/\|a_r\|_2$ never exceeds one: Flat-L2 enlarges no row, whereas Flat-GM

enlarges 22–32% of them, and the identity band and coefficient window leave 58–64% of the rows unscaled under Flat-L2 against 22–30% under Flat-GM. The generator information, in contrast, plays no role here. SA-Mat and Role-Hybrid produce byte-identical exports on all 306 models, and they differ from Flat-GM only on residual market-shortage rows of SG-Ter-Mer and Full that neither covers. Because the coupling rows of these models receive identical factors under Flat-GM, Role-Hybrid and SA-Mat, this identity does not test ownership; the ownership comparison requires the synthetic matrices of the next subsection.

7.2 A kernel effect mistaken for a block effect

In the archived heterogeneous headline cell, the median spectral diagnostic is 5.30×10^8 for Base, 1.36×10^4 for Flat-GM, 1.33×10^3 for SA-Mat and 4.68×10^2 for Flat-L2. The factor of about 10.2 between Flat-GM and SA-Mat was originally read as the value of block metadata. That comparison also changes the coupling kernel. With the kernel matched, no ownership component remains: Role-Hybrid, which uses no block scales, matches SA-Mat in all 30 archived records, and permuting the block-scale assignment leaves the diagnostic unchanged. [Proposition 4.2](#) gives the reason: after GM local normalization, the block-relative formula is Euclidean normalization. Two further tests agree. Pre-local block scales (SA-Pre) are worse than Role-Hybrid on all 30 instances, with median paired ratio 4.27, and in the exhaustive 24-permutation assignment sweep the true assignment is never the best (median rank 18).

The fresh replication reproduces this attribution on 180 new instances ([Figure 1](#)). All 1,080 exports pass the whole-row proportionality check at relative tolerance 10^{-13} . SA-Mat and Role-Hybrid are byte-identical on 136 of the 180 instances, and SA-Mat and its permuted-scale counterpart on 165. Where they differ, the implemented coefficient-window vetoes and filtered block extrema lie outside the hypotheses of [Proposition 4.1](#), and the effect is small: the SA-Mat/Role-Hybrid spectral ratio lies between 0.939 and 1.079 (median one), and the permutation ratio between 0.99962 and 1.00022. By contrast, the median SA-Mat/Flat-GM ratio falls from 0.937 to 0.102 and 0.00973 as the exponent range grows from $[-1, 1]$ to $[-3, 3]$. The large, heterogeneity-dependent advantage is therefore a coupling-kernel effect; ownership information accounts at most for the small, implementation-dependent differences in panel (b). Flat-L2 attains the smallest diagnostic in the archived cell because it also changes local normalization, and on the dispatch bases of [Section 7.1](#) the same kernel leaves conditioning essentially unchanged: kernel rankings on synthetic matrices need not transfer.

7.3 Completed-run gains are not failure-robust

On the 14 SG-Ter-Mer instances that every policy completed, Flat-GM used less solver work than Base, with paired geometric-mean ratios 0.698 at 1% gap and 0.682 at 0.1% (bootstrap intervals $[0.486, 0.944]$ and $[0.470, 0.925]$): about 30–32% less work on that pool. The fifteenth instance, instance077, completes under Base and times out under every scaled policy at both gaps. Under PAR10 this single failure dominates: Base averages 116 and 134 seconds, while every scaled policy averages between 1,216 and 1,218 seconds ([Online Resource 1, Table S3](#)). Scaling therefore produced a sizeable completed-run efficiency gain in this class, but the gain was not robust to failure-sensitive aggregation. The Simple class has no uniform ordering: Flat-GM lowers mean PAR10 from 11.9 to 2.1 seconds at 1% gap and raises it from 52.0 to 73.0 seconds at 0.1%. Role-Hybrid stays within one second of Flat-GM in all four cells.

The two endpoints answer different questions: completed-run effort describes how a representation affects the runs that finish, and PAR10 whether it changes which runs finish. With one seed and 15




Figure 1: Fresh matrix audit. Each exponent range contains 60 generated instances. Panel (a) changes the coupling kernel; panel (b) holds the kernels fixed and removes (Role-Hybrid) or permutes (SA-Mat-permB) the block-scale information. Marks are individual ratios with deterministic horizontal jitter; short horizontal segments are medians. The panels use different vertical scales. Unscaled matrices, whose estimates approach binary64 resolution, enter no ratio; all paired scaled diagnostics are below 6.5×10^6 .




Figure 2: Verified optima by coefficient multiplier. Each cell contains 48 configurations (three families, eight seeds, two presolve settings) per policy and solver. The fixed margin in (7) brings the continuous feasibility boundary increasingly close to an integer activity. Hatched cells are abstentions. Below the separator, GCD-floor is the later integer-aware reformulation control.

fixed instances per class, the evidence is a measured conditional gain, not a population estimate of acceleration. The same runs also show why a stated budget is needed to read residuals: an absolute residual of 4.97 on a zero-RHS equality of caso46e has an activity-normalized value of 1.15×10^{-7} , and without an application budget neither number establishes acceptance.

Table 4: Paired verification outcomes on identical model, presolve and solver settings (192 per comparison and solver). “*A* only” counts configurations verified under *A* and not under *B*.

Comparison <i>A</i> vs <i>B</i>	HiGHS		Gurobi		CPLEX	
	<i>A</i> only	<i>B</i> only	<i>A</i> only	<i>B</i> only	<i>A</i> only	<i>B</i> only
Budget-GM vs GM	41	0	12	0	0	0
Budget-dyadic vs Dyadic-GM	19	0	13	0	0	0
Budget-only vs GM	50	0	11	0	0	0
GM vs Base	0	26	0	16	0	10
Budget-only vs Base	24	0	1	6	0	10
Budget-GM vs Base	15	0	1	5	0	10
Budget-dyadic vs Base	2	8	1	4	0	10
GM-tight vs GM	25	0	26	0	26	0
Budget-only vs GM-tight	25	0	4	19	0	26
Budget-only vs Budget-GM	9	0	0	1	0	0

7.4 Budgets improve verification; acceptance is pipeline-dependent

The residual budget has a consistent effect within each scaling rule (Table 4). Plain GM centering loses verified optima relative to Base on all three solvers, 26 configurations on HiGHS, 16 on Gurobi and 10 on CPLEX, with no gains. At multipliers 10^6 and 10^9 it shrinks the fixed 0.1 margin of (7) below the solver tolerance, and most of its failures (49 of 51 on HiGHS, 36 of 52 on Gurobi, 34 of 52 on CPLEX) return points that satisfy the exported tolerance while violating the original row by more than the budget. Adding the budget to the same rule never lost a verified configuration on any solver. It recovered 41 of them on HiGHS and 12 on Gurobi, and none on CPLEX, where a rule and its budgeted version verify exactly the same configurations; the dyadic pair shows the same pattern, 19, 13 and 0.

Relative to the unscaled model, the outcome depends on the solver. On HiGHS, Budget-only verifies 191 of 192 configurations and Budget-GM 182, against 167 for Base and 141 for GM; Budget-only gains 24 configurations over Base and loses none. On Gurobi, Budget-only verifies 151, Budget-GM 152 and Base 156. On CPLEX every scaled policy verifies 140 and Base 150: the budget changes nothing there, and the ten configurations lost to scaling are lost with it and without it. On both commercial solvers, tightening the two tolerances to 10^{-9} (GM-tight, 166 configurations on each of the three solvers) is the more effective intervention. The same stored models and row factors thus lead to different acceptance in three pipelines, and the budget helps where the solver’s own handling of the exported coefficients is weakest.

The gap between Budget-only and Budget-GM on HiGHS is informative in its own right. At multipliers 10^6 and 10^9 , the only ones at which their outcomes differ, the budget endpoint is binding and Budget-GM selects that endpoint rounded inward, one unit in the last place (ulp) away from Budget-only’s factor. All ten discordant configurations, nine on HiGHS and one on Gurobi, with none on CPLEX, have exactly this one-ulp difference. The nine-configuration gap therefore measures how sensitive solver acceptance is to the last bit of a factor near a tolerance boundary, not an effect of coefficient centering. The same sensitivity appears at a larger scale: Budget-dyadic, which uses the nearest admissible power of two and exports every product exactly, verifies 161 configurations on HiGHS, 153 on Gurobi and 140 on CPLEX. Exact representation of the scaled row does not by itself secure acceptance.

Two controls delimit the result. GCD division with RHS flooring verifies all 192 configurations on each solver: it removes the deliberately narrow integer margin, whereas a budget controls the

original-unit meaning of an exported residual, so the budget result is a comparison among scalings and not an advantage over integer-aware presolve. Budget-round, which reserves the worst-case rounding term of (6), verifies all 48 configurations it accepts per solver and abstains on the other 144, including many that other policies solve correctly. Its abstentions certify that the requested uniform contract cannot be asserted; they are not solved instances, and Figure 2 counts them separately.

8 Discussion

8.1 What we learned

Row scaling can substantially change numerical conditioning. With the basis held fixed, the representation alone moves the condition number by six to eight orders of magnitude in the median, and the kernel decides whether it does.

The mechanism is simpler than the metadata explanation. Under exact GM local normalization, the extrema-based post-local block statistic carries no block-scale information (Proposition 4.1), so the block-relative rule becomes Euclidean normalization of the coupling rows, the optimal common factor of the idealized triangular model.

Matching kernels is essential for attribution. The original comparison changed the coupling kernel together with the information used, and the fixed-basis audit shows how much a kernel alone can change: identical coverage, representations orders of magnitude apart. A claim that generator information improves preprocessing needs a kernel-matched control.

Residual budgets are a meaningful invariant quantity. A budget carried with its row is unchanged by consistent re-emission, determines the admissible factors exactly, and in the stress test repaired the verification losses of coefficient centering on two of three solvers without introducing new ones on any of them. GCD division with RHS flooring verified every stress configuration, but it answers a different question: it is an integer-aware reformulation that exploits the activity lattice of integer rows, whereas a residual budget is a contract in original units that also applies to continuous rows and to constraints admitting no such reduction.

Solver acceptance and effort are properties of the pipeline. The same representation reduced completed-run work and caused a timeout, the same factors were accepted differently by the three solvers, and a one-ulp change in a factor changed HiGHS outcomes.

8.2 What remains limited

The operational evidence comes from one application, with two classes in the solver campaign, one solver, one seed and 15 fixed instances per class. The stress instances are small and deliberately adversarial; their budget of 0.01 is a numerical test parameter, and engineering budgets for the dispatch models have not been supplied or validated. Solver outcomes depend on the solver, its version and its settings. The collapse result concerns the extrema-based block statistic after GM normalization; other block statistics, other kernels and physical scale information are outside it, and on the operational models the coupling rows receive the same factors under every GM-local policy, so ownership is tested only on synthetic matrices. Finally, row scaling cannot repair within-row heterogeneity beyond the limit in (3) or the integer-rounding amplification in (6).

9 Conclusion

External row scaling genuinely matters numerically. Holding one basis fixed, it changes the conditioning of generated dispatch models by six to eight orders of magnitude in the median while preserving the MILP in exact arithmetic, and the kernel determines whether this happens.

Correct attribution changes the explanation of the gain. The advantage originally attributed to block metadata is a coupling-kernel effect: exact GM local normalization erases the extrema-based post-local block statistic, and a Euclidean coupling normalization, optimal in the idealized triangular model, accounts for the improvement without ownership information.

Residual budgets provide a constructive way to encode information that keeps its meaning in original application units. They come with exact feasibility and incompatibility conditions, a centered factor, invariance under re-emission, and explicit rounding and binary64 export conditions, and in the stress experiment they removed the verification losses caused by coefficient centering on two of the three solvers tested, without ever causing one.

What a solver does with a scaled representation still has to be measured, and its output still has to be verified. The practical lesson is therefore not to avoid external scaling, but to make it accountable: attribute its gains with matched controls, state its accuracy contract in the model’s own units, and verify solver output against that contract.

Data and software availability

The accompanying repository contains the C++ generator, frozen historical aggregates, the reconstructed operational logs and analysis, the fixed-basis audit with its exports, bases and independent verification, the exact-enumeration experiment, the row-factor implementation, tests, and the scripts that generate the tables and vector figures. The current research sources are distinguished from the historical draft, whose positive block-attribution claim is superseded. All newly reported data are stored under `results-revision/research-audit/`. No commercial solver is required to regenerate the figures, rerun the fixed-basis audit, verify the saved binary solutions, or run the HiGHS stress arm. Rerunning the Gurobi and CPLEX arms requires the corresponding licenses.

Statements and Declarations

Funding and competing interests. The author received no external funding for this work and declares no financial or non-financial competing interests.

Authorship and computational resources. Mathias Rodríguez Castro is the sole author. The computational platform originates in his hydrothermal-dispatch thesis work; the operational experiments used ClusterUY.

Use of AI assistance. OpenAI Codex was used to assist with mathematical derivations and their internal review, experiment and verification code, literature analysis, and drafting and revising the manuscript. Reported numerical results come from the stored computational runs. Figures were generated programmatically from those data or the stated equations. The author is responsible for the scientific content and the version submitted for publication.

Supplementary information

Online Resource 1 provides the triangular-model proof, the fixed-basis audit protocol and complete ranges, operational reconstruction details, frozen experiment protocols, arithmetic implementation checks, and reproduction instructions. The accompanying source and data archive is Online Resource 2.

References

- [1] Tobias Achterberg, Robert E. Bixby, Zonghao Gu, Edward Rothberg, and Dieter Weninger. Presolve reductions in mixed integer programming. *INFORMS Journal on Computing*, 32(2): 473–506, 2020. <https://doi.org/10.1287/ijoc.2018.0857>.
- [2] Timo Berthold and Gregor Hendel. Learning to scale mixed-integer programs. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(5):3661–3668, 2021. <https://doi.org/10.1609/aaai.v35i5.16482>.
- [3] Kevin K. H. Cheung, Ambros Gleixner, and Daniel E. Steffy. Verifying integer programming results. In Friedrich Eisenbrand and Jochen Koenemann, editors, *Integer Programming and Combinatorial Optimization*, volume 10328 of *Lecture Notes in Computer Science*, pages 148–160. Springer, 2017. https://doi.org/10.1007/978-3-319-59250-3_13.
- [4] A. R. Curtis and J. K. Reid. On the automatic scaling of matrices for Gaussian elimination. *IMA Journal of Applied Mathematics*, 10(1):118–124, 1972. <https://doi.org/10.1093/imamat/10.1.118>.
- [5] James Demmel. Nearly optimal Block-Jacobi preconditioning. *SIAM Journal on Matrix Analysis and Applications*, 44(1):408–413, 2023. <https://doi.org/10.1137/22M1504901>.
- [6] Joseph M. Elble and Nikolaos V. Sahinidis. Scaling linear optimization problems prior to application of the simplex method. *Computational Optimization and Applications*, 52(2): 345–371, 2012. <https://doi.org/10.1007/s10589-011-9420-4>.
- [7] GAMS Development Corporation. Special features for mathematical programs: Model scaling—the scale option. GAMS User’s Guide, online documentation, 2026. URL https://www.gams.com/latest/docs/UG_LanguageFeatures.html. Accessed September 8, 2026; documentation access year.
- [8] Gurobi Optimization, LLC. Tolerances and user-scaling. Gurobi Optimizer Reference Manual, online documentation, 2026. URL https://docs.gurobi.com/projects/optimizer/en/current/concepts/numericguide/tolerances_scaling.html. Accessed September 8, 2026; page updated June 17, 2026.
- [9] Nicholas J. Higham and Françoise Tisseur. A block algorithm for matrix 1-norm estimation, with an application to 1-norm pseudospectra. *SIAM Journal on Matrix Analysis and Applications*, 21(4):1185–1201, 2000. <https://doi.org/10.1137/S0895479899356080>.
- [10] Alexander Hoen and Ambros Gleixner. Analyzing the numerical correctness of branch-and-bound decisions for mixed-integer programming. In Guido Tack, editor, *Integration of Constraint Programming, Artificial Intelligence, and Operations Research*, volume 15763 of *Lecture Notes in Computer Science*, pages 35–50. Springer, 2025. https://doi.org/10.1007/978-3-031-95976-9_3.

- [11] Silke Horn. How do I scale my model to improve its numerical properties? Gurobi Help Center, 2025. URL <https://support.gurobi.com/hc/en-us/articles/26284047605521-How-do-I-scale-my-model-to-improve-its-numerical-properties>. Updated December 22, 2025; accessed September 8, 2026.
- [12] IDAES Project. CustomScalerBase class. IDAES v2.12.0 documentation, 2026. URL https://idaes-pse.readthedocs.io/en/stable/reference_guides/scaling/custom_scaler.html. Accessed September 8, 2026; documentation access year.
- [13] IDAES Project. Scaling methods. IDAES v2.12.0 documentation, 2026. URL https://idaes-pse.readthedocs.io/en/stable/reference_guides/core/util/scaling.html. Accessed September 8, 2026; documentation access year.
- [14] IEEE. *IEEE Standard for Floating-Point Arithmetic*. IEEE, 2019. IEEE Std 754-2019. <https://doi.org/10.1109/IEEESTD.2019.8766229>.
- [15] Jordan Jalving, Sungho Shin, and Victor M. Zavala. A graph-based modeling abstraction for optimization: Concepts and implementation in Plasmojl. *Mathematical Programming Computation*, 14:699–747, 2022. <https://doi.org/10.1007/s12532-022-00223-3>.
- [16] Ed Klotz. Identification, assessment, and correction of ill-conditioning and numerical instability in linear and integer programs. In *Bridging Data and Decisions*, INFORMS TutORials in Operations Research, pages 54–108. INFORMS, 2014. <https://doi.org/10.1287/educ.2014.0130>.
- [17] Robert Pritchett, Kathleen Howell, and Daniel Grebow. Low-thrust transfer design based on collocation techniques: Applications in the restricted three-body problem. In *AAS/AIAA Astrodynamics Specialist Conference*, Stevenson, Washington, 2017. URL https://engineering.purdue.edu/people/kathleen.howell.1/Publications/Conferences/2017_AAS_PriHowGre.pdf. AAS Paper 17-626.
- [18] Pyomo Development Team. ScaleModel: Model scaling transformation. Pyomo 6.9.1 documentation, 2026. URL <https://pyomo.readthedocs.io/en/6.9.1/api/pyomo.core.plugins.transform.scaling.ScaleModel.html>. Accessed September 8, 2026; documentation access year.
- [19] Zhaonan Qu, Wenzhi Gao, Oliver Hinder, Yinyu Ye, and Zhengyuan Zhou. Optimal diagonal preconditioning. *Operations Research*, 73(3):1479–1495, 2025. <https://doi.org/10.1287/opre.2022.0592>.