

Pre-Export Row Scaling for Generated Mixed-Integer Programs: Mechanism Attribution and Residual-Budget Contracts

Mathias Rodríguez Castro

Received: date / Accepted: date

Abstract Multiplying a row of a mixed-integer program by a positive number preserves it exactly and changes the numbers a solver sees. We ask what that change buys and what it still guarantees. On 306 generated hydrothermal-dispatch models, holding each model’s LP basis fixed across its representations, geometric-mean row scaling lowers its condition number on every model, by six to eight orders of magnitude in the median of each class. An apparent block-metadata effect is a coupling-kernel effect: geometric-mean centering makes the extrema-based block statistic identically one, so the block-relative rule becomes Euclidean normalization, which a control without ownership information reproduces; the idealized coupling problem has optimal factor $(1 + \rho^2)^{-1/2}$. Residual budgets stated in application units, by contrast, survive consistent re-emission, and we characterize exactly the row factors compatible with a budget and coefficient limits. In reconstructed dispatch runs with Gurobi, scaling used 30–32% less solver work on the 14 SG-Ter-Mer instances every policy completed, while one instance timed out under every scaled policy and reversed the failure-sensitive ordering, and reruns reproduce the gain with a second seed but not a second solver. In an exactly enumerated stress test, plain centering lost verified optima on all three solvers, adding the budget to that same rule lost none, and whether scaling helped at all differed across the three pipelines. A representation can be numerically better without being operationally safer, unless what safer means is stated in the units of the original model.

Keywords Mixed-integer linear programming · Row scaling · Numerical reliability · Equilibration · Residual budgets · Computational reproducibility

Mathematics Subject Classification (2020) 90C11 · 65F35 · 90-08 · 65G50

M. Rodríguez Castro
Facultad de Ingeniería, Universidad de la República, Montevideo, Uruguay
E-mail: mathiasr@fing.edu.uy

1 Introduction

A model generator can emit the same mixed-integer linear program (MILP) in many mathematically equivalent forms. Multiplying a constraint and its right-hand side (RHS) by a positive number preserves the MILP in exact arithmetic, yet it changes the coefficients on which presolve and LP factorizations operate and the original-unit meaning of an absolute residual tolerance. Solvers scale internally; a generator can also scale before export and use information that a matrix file does not carry, such as row roles, variable ownership and application units [11, 13].

The numerical effect of external row scaling can be large. On 306 dispatch models from a hydrothermal generator, holding one LP basis fixed across representations, geometric-mean row scaling lowers the 1-norm condition number of that basis on every model, by six to eight orders of magnitude in the median of each class (Section 7.1). A gain of this kind leaves one question open, in two halves:

when external preprocessing changes a generated MILP, what caused the numerical improvement, and what guarantee about the original model survives the change?

Both halves are about the same thing, the information the generator holds: which part of it explains the gain, and which part of it can be turned into a guarantee. The first half is answered by identifying the mechanism, the second by stating a contract in the units of the original model and checking it against what a solver returns. Representation quality, solver effort, failure-sensitive performance and verified accuracy are different endpoints, and evidence for one does not transfer to another.

This paper answers both halves for external row scaling of generated MILPs. Its contributions are the following.

- (A) *Attribution.* A comparison identifies the effect of block ownership only when coverage, row roles, kernels and safeguards are matched. With such controls, the order-of-magnitude conditioning advantage of a block-relative coupling rule over flat geometric-mean normalization is reproduced by a rule that uses no block scales, including its growth with block heterogeneity (Sections 4 and 7.2).
- (B) *Mechanism.* Geometric-mean row centering makes the post-local extrema-based block statistic identically one, so the block-relative coupling formula reduces to Euclidean row normalization. After local whitening, the idealized coupling problem has the closed-form optimal common factor $(1 + \rho^2)^{-1/2}$, which requires no ownership metadata (Section 4).
- (C) *Contract.* Residual budgets stated in application units survive consistent row re-emission. We characterize the admissible row factors exactly, with necessary and sufficient compatibility conditions, an incompatibility certificate, the constrained logarithmic minimax factor, the integer-rounding term that scaling cannot remove, and exact binary64 export with power-of-two factors (Section 5).

- (D) *Validation.* A fixed-basis audit of 306 dispatch models, a fresh audit of 1,080 matched matrix exports, a reconstruction of 240 operational solver runs with two matched sensitivity reruns of the same size, and an exactly enumerated stress experiment with 4,752 optimizer calls on three solvers (Sections 6 and 7).

Contributions (A) and (B) answer the first half of the question, (C) the second, and (D) is the evidence for both. The results progress from the representation to the solver. Scaling changes the numerical representation by orders of magnitude, and the kernel determines how much. The apparent metadata effect is a kernel effect with a simpler explanation. Residual budgets have measurable consequences: adding one to a centering rule never lost a verified stress configuration, relative to that same unbudgeted rule, on any of the three solvers tested. Whether a solver accepts the resulting points, and how much effort it spends, remain properties of the solver pipeline, to be measured and verified in original units.

2 Related work

Diagonal scaling has a long history, including the logarithmic method of Curtis and Reid [4]. Elble and Sahinidis [6] compare scalers for simplex and find that additional complexity need not improve performance. Klotz [16] relates numerical instability to formulation and finite precision. Optimal diagonal preconditioning is a broader problem: Qu et al. [19] use quasiconvex and semidefinite formulations, while Demmel [5] gives block-Jacobi guarantees. Our triangular calculation is a scalar special case of this problem that explains the attribution result.

Robustness is also an established objective in MIP scaling. Berthold and Hendel [2] learn to select between standard and Curtis–Reid scaling in Xpress from an indicator of numerical difficulties. Their implementation uses power-of-two factors and their evaluation examines inconsistencies across seeds. Our external construction instead states row-level admissibility conditions; it does not predict or replace internal solver scaling.

Modeling environments already support generator-level information: Pyomo’s `ScaleModel` applies component factors [18], GAMS has variable and equation scale attributes [7], and IDAES derives and propagates scales across submodels [12, 13]. Plasmol retains local problems and linking constraints [15]. Consequently, using modeling metadata before export is not new. Our identifiability result concerns the loss of one specific block statistic after a preceding normalization.

Tolerance-based scaling also has direct precedent. Gurobi documents absolute feasibility and integrality tolerances and the treatment of very small coefficients [8], and recommends external scaling that makes tolerance-sized violations negligible in application units [11]. Pritchett, Howell and Grebow [17] similarly scale constraints using physical accuracy requirements and SNOPT’s

tolerance. Section 5 combines this known transfer with coefficient limits, identifies incompatible specifications, and solves the constrained centering problem. Integer rounding and exact binary64 export impose additional restrictions.

Integer-aware presolve can instead strengthen the formulation. GCD division and RHS rounding are established reductions [1]; on integer rows they can remove a narrow feasibility margin that pure scaling preserves. We include that distinction in the stress comparison, rather than treating tolerance-based scaling as a replacement for presolve.

Finally, Cheung, Gleixner and Steffy [3] provide certificates for integer-programming results, and Hoen and Gleixner [10] analyze individual branch-and-bound decisions. Our rational residual checks address primal acceptance, and independent enumeration supplies optimality on the small binary instances. What this paper adds to these lines of work is attribution: matched comparisons of actual exports that separate what a scaling policy changes from which information causes the change.

3 Model, policies and comparisons

Consider

$$\min c^T x \quad \text{subject to } Ax \leq b, \quad l \leq x \leq u, \quad x_j \in \mathbb{Z} \ (j \in J),$$

where equalities and opposite inequality senses are allowed. For a positive diagonal matrix D_r , replacing (A, b) by $(D_r A, D_r b)$ preserves every constraint, variable domain and objective value in exact arithmetic. The whole row, including its RHS, is multiplied by one positive number; dividing different coefficients of a row by different block scales is not a row scaling and is outside this equivalence.

The generator provides a partition of the local rows into nonempty blocks, a set of coupling rows, and a column-to-block map β . A scaling policy is then determined by four separate choices: *coverage*, the rows eligible for scaling; the *kernel*, which maps an eligible row to a candidate factor; *safeguards*, which decide whether a candidate is applied; and *information*, the generator data the kernel may read, such as roles, block scales and ownership. An effect can be attributed to information only when the other three choices are matched.

For a nonzero row, put $m_r = \min_{a_{rj} \neq 0} |a_{rj}|$ and $M_r = \max_j |a_{rj}|$. The geometric-mean (GM) factor is $(m_r M_r)^{-1/2}$ and the Euclidean (L2) factor is $1/\|a_r\|_2$. Both are coefficient-only, but they are different kernels, and Table 1 keeps the distinction visible.

Relabeling blocks without changing their assignment is an invariance check, whereas permuting the assignment of unequal block scales tests whether the map is useful on the chosen instances.

Table 1 Policies. Coverage and safeguards are matched across the comparisons in which a policy is used; a role map and a block-scale map are separate information sources. The block scales s_i are defined in Section 4.

Policy	Local rows	Coupling rows	Information used
Base	None	None	None
Flat-GM	GM	GM	Row coefficients
Flat-L2	L2	L2	Row coefficients
Role-Hybrid	GM	L2	Row roles, coefficients
SA-Mat	GM	$\left[\sum_j (a_{rj}/s_{\beta(j)})^2\right]^{-1/2}$	Roles, block scales, ownership
SA-Pre	GM	Same expression, pre-local scales	Roles, original block scales

4 Identifying the mechanism

This section shows that, under exact geometric-mean local normalization, the extrema-based post-local block statistic used by the original policy carries no block-scale information, so the coupling stage cannot use it, and it derives the simpler mechanism that remains.

4.1 Collapse of post-normalization block centers

For each nonzero local row r , write m_r and M_r for its smallest and largest nonzero coefficient magnitudes. The geometric-mean kernel applies $d_r = (m_r M_r)^{-1/2}$. Let $\tilde{m}_r = d_r m_r$ and $\tilde{M}_r = d_r M_r$ denote the exported extrema. For a nonempty set I of local rows, the block statistic is

$$s_I = \sqrt{\left(\min_{r \in I} \tilde{m}_r\right) \left(\max_{r \in I} \tilde{M}_r\right)}.$$

Proposition 4.1 (Exact collapse and the effect of a skip band) *If every row receives its geometric-mean factor, then $s_I = 1$ for every nonempty grouping I of local rows. More generally, suppose the only safeguard replaces d_r by 1 when the candidate lies in $(1/u, u)$, where $u \geq 1$. Then $s_I \in [1/u, u]$, independently of the original coefficient scales and of the grouping of the rows.*

Proof Under exact normalization, $\tilde{m}_r = \sqrt{m_r/M_r} = 1/\tilde{M}_r$. Consequently $\min_r \tilde{m}_r = 1/\max_r \tilde{M}_r$, proving the first statement. For the second, let $c_r = \sqrt{\tilde{m}_r \tilde{M}_r}$. An accepted row has $c_r = 1$; a skipped row has $c_r = \sqrt{m_r M_r} \in (1/u, u)$. Thus $u^{-2} \leq \tilde{m}_r \tilde{M}_r \leq u^2$ for every row. Choose p attaining the maximum exported row magnitude and q attaining the minimum. Then

$$u^{-2} \leq \tilde{m}_q \tilde{M}_q \leq \tilde{m}_q \tilde{M}_p \leq \tilde{m}_p \tilde{M}_p \leq u^2.$$

Taking square roots gives the result.

For a coupling row a_r , let $\beta(j)$ assign column j to a local block, and define the block-weighted candidate and its ownership-blind counterpart by

$$\gamma_r^{\text{block}} = \left(\sum_j a_{rj}^2 / s_{\beta(j)}^2 \right)^{-1/2}, \quad \gamma_r^{\text{L2}} = \left(\sum_j a_{rj}^2 \right)^{-1/2}.$$

Blockless columns may use an arithmetic mean of valid block scales or the unit fallback; both conventions preserve the interval in Proposition 4.1.

Proposition 4.2 (Loss of ownership information) *Under exact local normalization and with no coefficient filtering, $\gamma_r^{\text{block}} = \gamma_r^{\text{L2}}$ for every coupling row, so changing the ownership map cannot change the candidate export; applying the same coupling safeguards to identical candidates preserves equality. Under the skip-band assumptions of Proposition 4.1,*

$$u^{-1} \leq \frac{\gamma_r^{\text{block}}}{\gamma_r^{\text{L2}}} \leq u,$$

and, before coupling safeguards,

$$u^{-2} \leq \frac{\kappa_2^+(A_{\text{block}})}{\kappa_2^+(A_{\text{L2}})} \leq u^2,$$

where $\kappa_2^+(A) = \sigma_{\max}(A) / \sigma_{\min}^+(A)$ is taken over the positive singular values, so the bound covers rectangular and rank-deficient exports.

Proof The equality follows by substituting $s_i = 1$. In the second case,

$$u^{-2} \|a_r\|_2^2 \leq \sum_j a_{rj}^2 / s_{\beta(j)}^2 \leq u^2 \|a_r\|_2^2,$$

which gives the factor bound. Write the block-based export as EA_{L2} , where E has unit entries on local rows and entries in $[1/u, u]$ on coupling rows. Multiplication by E changes each singular value by a factor between $1/u$ and u and preserves rank, so $\sigma_{\max}$ and $\sigma_{\min}^+$ each move by at most u ; taking the two extremes in opposite directions gives the stated interval for their ratio.

Both statements assume that the block summary sees every coefficient. Coefficient-window vetoes and filtered extrema lie outside the hypotheses: a row with magnitudes $(10^{-8}, 3, 10^8)$ is geometrically centered, but restricting the summary to $(10^{-6}, 10^6)$ produces $s_I = 3$, a nonunit statistic created by the filter rather than by retained physical scale. The export audit in Section 7 measures how often the implemented safeguards break exact equality.

4.2 What a matched comparison identifies

With equal coverage and equal local kernels, two exports agree on their local rows, so any difference is confined to the coupling rows. This localizes the difference without attributing it to ownership, because the coupling kernels may still differ. The ownership control therefore keeps local GM and coupling L2 normalization but replaces every block scale by one (Role-Hybrid in Table 1); SA-Mat versus Role-Hybrid is the ownership comparison. Flat-GM changes the coupling kernel and Flat-L2 changes the local kernel, so neither isolates ownership.

Equality also requires equal *applied* factors, not merely an inactive coupling stage in one policy. For local rows $x_1 \leq 1$, $x_2 \leq 1$ and the coupling row $x_1 + 10^{-2}x_2 \leq 1$, both block scales are one and the Euclidean coupling candidate $(1 + 10^{-4})^{-1/2}$ falls inside the skip band $(1/2, 2)$, whereas the flat GM candidate 10 is applied. Both policies cover every row, yet their exports differ.

4.3 The spectral slope in a stylized model

Take $B \geq 2$ local singleton rows with positive coefficients $v_i = 10^{e_i}$, where $\min_i e_i = -R$ and $\max_i e_i = R$, and one coupling row $v^\top$. Exact local normalization produces I_B and unit post-local block centers. A whole-row coupling factor γ gives

$$A_\gamma = \begin{bmatrix} I_B \\ \gamma v^\top \end{bmatrix}, \quad A_\gamma^\top A_\gamma = I_B + \gamma^2 v v^\top, \quad \kappa_2(A_\gamma) = \sqrt{1 + \gamma^2 \|v\|_2^2}.$$

Flat geometric-mean normalization selects $\gamma = 1$, whereas both the post-local block rule and the hybrid Euclidean control select $\gamma = 1/\|v\|_2$. Hence

$$\kappa_2(A_{\text{GM}}) = \sqrt{1 + \sum_i 10^{2e_i}}, \quad \kappa_2(A_{\text{block}}) = \kappa_2(A_{\text{hybrid}}) = \sqrt{2}.$$

For fixed B the ratio is of order 10^{-R} : the advantage grows with block heterogeneity and survives removal of the ownership map. Every coupling coefficient is multiplied by the same γ , so all within-row ratios are preserved; dividing each v_i by its own 10^{e_i} would change the constraint and is not a row scaling.

4.4 Optimal coupling factor after local whitening

The stylized model suggests a constructive replacement for the metadata explanation: once local scale has been removed, choose the coupling factor from the coupling coefficients themselves. In the idealized triangular model this choice has an exact solution.

Proposition 4.3 (Optimal common coupling factor after local whitening) *Let $R \in \mathbb{R}^{m \times n}$, with $m, n \geq 1$, and set $\rho = \|R\|_2$. For*

$$C_\gamma = \begin{bmatrix} I_n & 0 \\ \gamma R & \gamma I_m \end{bmatrix}, \quad \gamma > 0,$$

the unique factor minimizing $\kappa_2(C_\gamma)$ is

$$\gamma_* = (1 + \rho^2)^{-1/2}, \quad \min_{\gamma > 0} \kappa_2(C_\gamma) = \sqrt{1 + \rho^2} + \rho.$$

On an admissible interval $0 < \ell \leq \gamma \leq h$, the optimum is $\min\{h, \max\{\ell, \gamma_\}\}$.*

The proof (Online Resource 1) reduces the spectrum by an SVD to two-by-two blocks. For the block at ρ , the singular-value ratio k satisfies $k + k^{-1} = \gamma^{-1} + \gamma(1 + \rho^2)$, whose unique minimizer gives the result.

For a single coupling row, γ_* is the Euclidean normalization of the complete row, including the coefficient of its aggregate variable. The idealized coupling problem therefore has a simple optimal whole-row normalization that uses no ownership information, and this is the mechanism the matched control isolates empirically. The statement concerns a common factor on a locally whitened triangular structure; for general generated MIPs it motivates the Euclidean coupling kernel rather than certifying it.

Why conditioning needs a residual bound. The interval form of Proposition 4.3 matters. In A_γ , κ_2 decreases to one as $\gamma \rightarrow 0$, while a fixed residual on the exported coupling row corresponds to an original residual larger by the factor $1/\gamma$. A conditioning criterion alone cannot say how small a coupling factor may be: on any interval its minimizer for A_γ is the lower endpoint ℓ . Section 5 derives that endpoint from an application-unit residual budget. The two halves of the paper meet here. In this construction, block ownership did not survive the preceding normalization as usable scaling information; application-unit accuracy does survive it, once it is encoded as a residual budget carried with the row.

5 Residual-budget contracts

The first half of the question asked which generator information explains the numerical gain, and the answer was that block ownership does not. This half asks which generator information can be turned into a guarantee. Ownership is one kind of generator information; a tolerance stated in application units is another. A residual budget says how much violation of an original constraint is acceptable. It must be supplied by the application or declared as a numerical test parameter; it cannot be inferred from a block label or a zero right-hand side. Scaling rows so that solver tolerances correspond to negligible application-unit violations is established practice [8, 11, 13]. The results below make its combination with coefficient limits exact: they characterize every admissible

factor, certify when none exists, select the centered one, and identify what row scaling cannot control. Individually the steps are elementary algebra; what the section contributes is that the contract is complete, in the sense that a factor is admissible exactly when the conditions hold, and that each consequence is stated in the units the application uses.

For a nonconstant row $a^T x \leq b$, define the original residual $v(x) = (a^T x - b)_+$ and a budget $\delta > 0$. Let $\epsilon > 0$ bound the residual of the *exported* row. This is a hypothesis about the returned point rather than a consequence of a solver parameter: presolve, internal scaling and finite precision intervene between the two. For a positive row factor α ,

$$(\alpha a^T x - \alpha b)_+ = \alpha v(x) \quad (1)$$

in exact arithmetic, so an exported residual at most ϵ implies $v(x) \leq \delta$ whenever $\alpha \geq \epsilon/\delta$. Absolute residuals give the same identity for equalities. Write $m = \min_{a_j \neq 0} |a_j|$ and $M = \max_j |a_j|$, and let $0 < \ell \leq U$ limit the exported coefficient magnitudes. An additional RHS limit U_b intersects the upper endpoint below with $U_b/|b|$ for $b \neq 0$.

Proposition 5.1 (Admissible factors and constrained centering) *The factors that place every nonzero coefficient in $[\ell, U]$ and satisfy $\epsilon/\alpha \leq \delta$ form exactly the interval*

$$\mathcal{I} = [L, H], \quad L = \max\{\epsilon/\delta, \ell/m\}, \quad H = U/M. \quad (2)$$

The interval is nonempty if and only if

$$M/m \leq U/\ell \quad \text{and} \quad \delta \geq \epsilon M/U. \quad (3)$$

When it is nonempty, the unique minimizer of

$$\min_{\alpha \in \mathcal{I}} \max_{a_j \neq 0} |\log(\alpha |a_j|)| \quad (4)$$

is

$$\alpha^* = \min\{H, \max\{L, (mM)^{-1/2}\}\}. \quad (5)$$

Proof The coefficient bounds are equivalent to $\alpha m \geq \ell$ and $\alpha M \leq U$. Together with $\alpha \delta \geq \epsilon$ they give (2). The condition $L \leq H$ is equivalent to the two inequalities in (3). To solve the objective, put $t = \log \alpha$, $q = (\log m + \log M)/2$ and $h = (\log M - \log m)/2$. Since the extreme absolute logarithms occur at an endpoint,

$$\max_j |t + \log |a_j|| = \max\{|t + \log m|, |t + \log M|\} = h + |t + q|.$$

The unique minimizer on $[\log L, \log H]$ is the projection of $-q$ onto that interval. Exponentiating gives (5).

Incompatibility certificate. When $L > H$, (3) names the violated requirement. A within-row spread $M/m > U/\ell$ cannot be repaired by any row factor; a budget $\delta < \epsilon M/U$ conflicts with the coefficient ceiling even for a perfectly balanced row. In both cases the certificate concerns the numerical specification, not the optimization model: exporting the unscaled row does not meet the contract, and resolving the conflict requires revising units, limits or the formulation. The factor α^* optimizes coefficient centering subject to the contract; it is not a condition-number minimizer.

Corollary 5.2 (Invariance under consistent re-emission) *Replace (a, b, δ) by $(qa, qb, q\delta)$ for any $q > 0$, keeping ϵ, ℓ, U fixed. The interval becomes $\mathcal{I}/q$, abstention is unchanged, and the selected factor becomes α^*/q . The exported row and the budget-normalized original residual $v(x)/\delta$ are therefore unchanged.*

Proof Both coefficient extremes and the original budget acquire the factor q . Each term defining L , H and $(mM)^{-1/2}$ is divided by q . Equation (1) gives the residual statement.

A budget is therefore information that survives preprocessing: generators that emit the same constraint in different units, each with the budget expressed in its own units, obtain the same exported row and the same contract. The invariance is lost if a row is re-emitted without transforming its budget, and a fixed cutoff on the factor itself does not have it. Restricting factors to powers of two [2] retains it for power-of-two re-emissions, subject to the representability conditions of Proposition 5.3.

5.1 The integer-rounding term

Let J be the integer columns and let $\bar{x}$ be obtained by rounding only these coordinates of x . Suppose $|\bar{x}_j - x_j| \leq \eta_j$ for $j \in J$ and put $E = \sum_{j \in J} |a_j| \eta_j$. The triangle inequality and (1) give

$$v(\bar{x}) \leq \epsilon/\alpha + E, \quad (6)$$

also for equalities with absolute residuals. Changing the row factor reduces the exported-residual term but leaves the rounding term, which is measured in original units, unchanged. Normalizing a big- M row therefore cannot remove the effect of an integrality tolerance on that row.

For a contract on the rounded assignment, Proposition 5.1 applies with δ replaced by $\delta - E$ when $\delta > E$; if $\delta \leq E$, no positive factor certifies it. The bound is a worst case: actual rounding errors may be smaller or cancel, and other rows can exclude the worst case, so an empty interval shows that the stated information is insufficient rather than that returned points will violate the budget. Under re-emission E becomes qE , so invariance is retained. A tighter integrality tolerance or a reformulation can reduce E ; row scaling cannot.

5.2 Exact export on a dyadic grid

A factor in the real admissible interval does not guarantee an exactly proportional exported row: the products $\text{fl}(\alpha a_j)$ and $\text{fl}(\alpha b)$ are rounded independently, and rounding a computed endpoint to nearest can leave the closed interval. Our arbitrary-factor implementation forms the endpoints as exact rational numbers in the supplied binary64 inputs, rounds them *inward*, and abstains if no positive binary64 factor remains. This certifies admissibility of the factor, not exactness of the products.

Powers of two avoid multiplication rounding. The following intersection also accounts for subnormal numbers and for the RHS. Assume finite binary64 input data and IEEE arithmetic with gradual underflow [14]. Write each nonzero datum $z \in \{a_1, \dots, a_n, b\}$ uniquely as $z = \text{sign}(z)N_z 2^{p_z}$ with N_z a positive odd integer, and put $q_z = \lfloor \log_2 |z| \rfloor$. Let $[L, H]$ be the desired real factor interval, including any rounding allowance from (6).

Proposition 5.3 (Exact dyadic admissibility) *A factor $\alpha = 2^k$ is positive and representable in binary64, belongs to $[L, H]$, and multiplies every nonzero coefficient and the RHS exactly into a finite binary64 value if and only if the integer k satisfies*

$$K_- = \max \left\{ -1074, \lceil \log_2 L \rceil, \max_z (-1074 - p_z) \right\} \leq k,$$

$$k \leq K_+ = \min \left\{ 1023, \lfloor \log_2 H \rfloor, \min_z (1023 - q_z) \right\}.$$

If $K_- > K_+$ there is no such factor. Otherwise an exponent nearest to $-\frac{1}{2} \log_2(mM)$ among the integers in $[K_-, K_+]$ minimizes the logarithmic coefficient objective (4) on this exact-export set; a tie may be resolved by choosing the larger factor.

Proof The factor itself requires $-1074 \leq k \leq 1023$. Since the datum's odd significand is unchanged, the product $N_z 2^{p_z+k}$ lies on the subnormal grid exactly when $p_z + k \geq -1074$; its highest bit stays below overflow exactly when $q_z + k \leq 1023$. The original significand has at most 53 bits, so these two conditions are also sufficient. Intersecting them over the row data and with $L \leq 2^k \leq H$ gives the interval of integer exponents. The objective is $\frac{1}{2} \log(M/m) + |k \log 2 + \frac{1}{2} \log(mM)|$, proving the nearest-exponent claim.

The implementation obtains the endpoint exponents from integer ratios and bit lengths, compares candidate objectives exactly as rational numbers, and checks the rounded products against exact rational products. Writing the row to a text file additionally requires a round-trip decimal serializer, since fixed decimal places can discard small values even when the multiplication is exact.

Table 2 Endpoints and how they are measured.

Endpoint	Measurement	Results
Exact equivalence	Whole-row proportionality of named exports, including RHS	§7.1, §7.2
Numerical representation	κ_1 of one fixed basis; κ_2^+ of exported matrices	§7.1, §7.2
Mechanism attribution	Matched exports: byte identity and paired spectral ratios	§7.2
Solver effort	Gurobi work units on runs completed by every policy	§7.3
Failure-sensitive performance	PAR10 over a fixed instance pool	§7.3
Solver acceptance	Solver status and returned point	§7.4
Verified correctness	Exact original residuals, bounds and integrality; enumerated optimum	§7.4

Verification after solving. The contract is conditional on the exported residual, so returned points are checked against the original rows, bounds and integrality restrictions. On the stress instances of Section 6 the check is exact: row residuals are rational expressions in the stored binary64 data, and optima are obtained by enumeration. On the dispatch models the archived checker uses floating-point residuals, which are diagnostics; certifying branch-and-bound results is a stronger task [3].

6 Computational design

The experiments measure the endpoints in Table 2. Each result is stated for its own endpoint and none is used as a proxy for another: a condition number describes a representation, not solver effort or the correctness of a returned point, and a solver status is not an independently verified solution.

6.1 Fixed-basis conditioning audit

A condition number read along each policy’s own solve compares different vertices as well as different representations. This audit holds the vertex fixed. For each of the 306 archived dispatch inputs (204 Simple, 68 SG-Ter-Mer and 34 Full models), the C++ generator exports Base, Flat-GM, Flat-L2, Role-Hybrid and SA-Mat without calling an optimizer. HiGHS 1.15.1 solves the Base LP relaxation with the simplex method on one thread, and its optimal basis is recorded by row and column names. The same basis is then assembled from every policy’s own export: structural basic columns come from the scaled matrix and basic slacks enter as unit columns, as a solver holds them. Before assembly, each export is checked to be a whole-row positive scaling of Base, including RHS values, with identical bounds, objective and integrality (relative tolerance 2×10^{-13}).

The diagnostic is $\kappa_1(B) = \|B\|_1 \|B^{-1}\|_1$, with $\|B^{-1}\|_1$ estimated from a sparse LU factorization by the block 1-norm estimator [9]. Because κ_1 is not invariant under row scaling, the audit measures the conditioning of the exported representation before any solver-internal scaling, not an intrinsic property of the vertex. An independent script rebuilds every basis from the LP text and recomputes each estimate through an exact power-of-two row equilibration, and a precision check bounds κ_1 from both sides on a stress sample. The instance set is complete, and all exports, bases and hashes are archived.

6.2 Matrix attribution audit

A C++ generator of block-structured instances builds local blocks and coupling rows and retains row roles and column ownership, so every policy in Table 1 is available in one implementation and exports can be compared before any optimizer is called. The archived headline cell has four heterogeneous blocks, coefficient exponent range $[-2, 2]$, coupling disturbance severity 6 and 30 seeds. A fresh replication crosses ten seeds (20260909–20260918) with exponent ranges $[-R, R]$ for $R \in \{1, 2, 3\}$, two severities and three coupling patterns: 180 instances and 1,080 exports, written with a serializer that preserves binary64 values on readback. Identity claims rest on direct comparisons of named exports. The spectral pseudo-condition number $\kappa_2^+(A) = \sigma_{\max}(A)/\sigma_{\min}^+(A)$ uses a numerical rank fixed per instance from the row-L2-normalized Base (relative threshold 10^{-12}) and shared by all its variants.

6.3 Operational reconstruction

The operational campaign runs Base, Flat-GM, Flat-L2 and Role-Hybrid on 15 fixed instances from each of the Simple and SG-Ter-Mer classes with Gurobi at relative MIP gaps of 1% and 0.1% (seed 42, 16 threads, 1,800-second limit): 240 assigned runs, whose native logs and input hashes were retrieved from the cluster. The subsets were fixed before outcomes were inspected. Nine of the 15 SG-Ter-Mer instances, including instance077, carry a scenario defect that makes them unfaithful market cases without affecting their validity as MILPs or the comparison of representations of one model; Online Resource 1 repeats the endpoints on the six instances without the defect. The pre-specified primary endpoint is failure-sensitive performance: the arithmetic-mean PAR10 over all 15 instances of a cell, with 18,000 seconds assigned to a non-completion. Solver effort, the secondary endpoint, is Gurobi’s work measure, recovered from the native logs at a common resolution of 0.01 units, on the instances that every policy completed; paired geometric-mean ratios carry exploratory bootstrap intervals (Online Resource 1). The two gap settings reuse the same instances and are not independent replications.

6.4 Enumerated residual-budget stress experiment

Three binary families with 12 variables contain a cardinality row, three positive packing rows and three rows with signed integer coefficients; profits are positive integers. With integer matrix W and integer capacities k , the represented constraints are

$$10^e Wx \leq 10^e k - 0.1 \mathbf{1}, \quad x \in \{0, 1\}^{12}, \quad e \in \{0, 3, 6, 9\}. \quad (7)$$

For these exponents the binary feasible set is exactly $Wx \leq k - \mathbf{1}$, and the stored binary64 right-hand sides are checked to separate the adjacent scaled integer activities. The exponent therefore changes the representation and the continuous relaxation without changing the binary optimum. Because the absolute margin 0.1 is fixed, this is not pure re-emission: it deliberately places infeasible integer points at a controlled distance from the continuous boundary. The budget is 0.01 in the units of the represented row, a declared test parameter.

Eight seeds (100–107), reserved after a one-seed pilot, give 96 models. Each is solved with presolve on and off by HiGHS 1.15.1, Gurobi 13.0.2 and CPLEX 22.1.2 on one thread, with zero target gap, a ten-second limit, and primal and integrality tolerances 10^{-6} . Base, GM and L2 scale every row by their kernels without an identity band; GM-tight uses GM with both tolerances tightened to 10^{-9} ; Budget-only uses the direct factor ϵ/δ_r with $\epsilon = 10^{-6}$; Budget-GM applies Proposition 5.1 with coefficient limits $[10^{-9}, 10^9]$; Budget-round additionally reserves $E_r = 10^{-6} \sum_j |a_{rj}|$ and abstains on the whole model if any row’s interval is empty. Dyadic-GM and Budget-dyadic use power-of-two factors under Proposition 5.3, without and with the budget. These two arms were added after an endpoint rounding issue was found in the initial run and are exploratory; the original seven arms were rerun with the corrected implementation and unchanged seeds. The nine policies give 1,728 planned configurations per solver; Budget-round abstains on 144 of them, leaving 1,584 calls per solver and 4,752 in total.

A configuration is *verified* when the returned point is within 10^{-9} of a binary assignment that is exactly feasible for the stored binary64 data and attains the optimum found by enumerating all 2^{12} assignments, and the bounds hold within 10^{-9} . Residual budgets and integrality errors are recorded separately, and the rule does not use the solver status. The instances are small and deliberately close to integer feasibility boundaries, so counts describe a failure mechanism rather than practical failure rates, and repeated solver settings on one model are not treated as independent samples. A later control on the same 96 stored models applies GCD division with RHS flooring, an established integer-aware reduction that preserves every binary feasible set but changes the LP relaxation; Online Resource 1 reports it together with two GM variants.

Table 3 Fixed-basis conditioning audit on all 306 dispatch models. One optimal basis of the Base LP relaxation is assembled from every policy’s own export. The first row is the median κ_1 of the Base basis; the others give the median per-instance ratio $\kappa_1(\text{policy})/\kappa_1(\text{Base})$ and, in parentheses, the number of instances with a ratio below one. Ranges are in Online Resource 1, Table S1.

Policy	Simple ($n = 204$)	SG-Ter-Mer ($n = 68$)	Full ($n = 34$)
Base, median κ_1	8.8×10^{18}	2.9×10^{20}	4.3×10^{20}
Flat-GM	2.4×10^{-7} (204)	7.3×10^{-9} (68)	3.5×10^{-7} (34)
Flat-L2	1.57 (32)	1.51 (10)	1.49 (2)
Role-Hybrid	2.4×10^{-7} (204)	7.3×10^{-9} (68)	3.5×10^{-7} (34)
SA-Mat	2.4×10^{-7} (204)	7.3×10^{-9} (68)	3.5×10^{-7} (34)

7 Results

7.1 Scaling materially changes numerical representation

All 1,530 exports pass the whole-row scaling check, with largest relative deviation 2.5×10^{-16} , and every policy is evaluated on the same basis. The three policies with GM local normalization lower κ_1 on every one of the 306 models, with median ratios of 2.4×10^{-7} in Simple, 7.3×10^{-9} in SG-Ter-Mer and 3.5×10^{-7} in Full (Table 3). Because the vertex is common, these are differences between representations rather than between the vertices that different solves reach. The independent reconstruction reproduces all 1,530 bases exactly. Because the unscaled values exceed the reciprocal of machine precision, their accuracy was also checked directly on a stress sample of 24 bases from eight models: in each class, the model with the largest Base estimate, the model with the weakest Flat-GM ratio and one seeded random model, under Base, Flat-GM and Flat-L2. A lower bound certified with exact rational residuals and an a posteriori upper bound over all columns of B^{-1} enclose every κ_1 within a relative width of 6×10^{-3} , and the resulting upper bounds on the Flat-GM/Base ratio agree with the estimated ratios to 2.1×10^{-5} (Online Resource 1, Table S2). The median separation of six to eight orders of magnitude, present on every model, is therefore a property of the representations and not of floating-point error in the unscaled bases. An archived computation on a different reference basis, the LP with integers fixed at a Gurobi incumbent, points in the same direction with smaller magnitudes, but its exports were not retained and it is not used as evidence (Online Resource 1).

The kernel, with the safeguards acting on its candidates, determines this separation. Flat-L2 covers the same rows as Flat-GM, yet its median ratio stays between 1.49 and 1.57 and it lowers κ_1 on only 44 of the 306 models. Every row of these models has a coefficient of magnitude at least one, so the Euclidean factor $1/\|a_r\|_2$ never exceeds one: Flat-L2 enlarges no row, whereas Flat-GM enlarges 22–32% of them, and the identity band and coefficient window leave 58–64% of the rows unscaled under Flat-L2 against 22–30% under Flat-GM. The generator information, in contrast, plays no role here. SA-Mat and Role-

Hybrid produce byte-identical exports on all 306 models, and they differ from Flat-GM only on residual market-shortage rows of SG-Ter-Mer and Full that neither covers. Because the coupling rows of these models receive identical factors under Flat-GM, Role-Hybrid and SA-Mat, this identity does not test ownership; the ownership comparison requires the synthetic matrices of the next subsection.

Flat-L2 is worth pausing on, because it is the same kernel that the matched synthetic control will show to be the strong one. On generated dispatch bases it is the weak policy, better than Base on 44 of 306 models; on the synthetic coupling matrices it reproduces the entire advantage attributed to block metadata. A kernel ranking obtained from one matrix diagnostic therefore does not transfer to another family, still less to solver behavior, and the endpoints of Table 2 have to be measured where they are claimed.

7.2 A kernel effect mistaken for a block effect

In the archived heterogeneous headline cell, the median spectral diagnostic is 5.30×10^8 for Base, 1.36×10^4 for Flat-GM, 1.33×10^3 for SA-Mat and 4.68×10^2 for Flat-L2. The factor of about 10.2 between Flat-GM and SA-Mat was originally read as the value of block metadata. That comparison also changes the coupling kernel. With the kernel matched, no ownership component remains: Role-Hybrid, which uses no block scales, matches SA-Mat in all 30 archived records, and permuting the block-scale assignment leaves the diagnostic unchanged. Proposition 4.2 gives the reason: after GM local normalization, the block-relative formula is Euclidean normalization. Two further tests agree. Pre-local block scales (SA-Pre) are worse than Role-Hybrid on all 30 instances, with median paired ratio 4.27, and in the exhaustive 24-permutation assignment sweep the true assignment is never the best (median rank 18).

The fresh replication reproduces this attribution on 180 new instances (Figure 1). All 1,080 exports pass the whole-row proportionality check at relative tolerance 10^{-13} . SA-Mat and Role-Hybrid are byte-identical on 136 of the 180 instances, and SA-Mat and its permuted-scale counterpart on 165. Where they differ, the implemented coefficient-window vetoes and filtered block extrema lie outside the hypotheses of Proposition 4.1, and the effect is small: the SA-Mat/Role-Hybrid spectral ratio lies between 0.939 and 1.079 (median one), and the permutation ratio between 0.99962 and 1.00022. By contrast, the median SA-Mat/Flat-GM ratio falls from 0.937 to 0.102 and 0.00973 as the exponent range grows from $[-1, 1]$ to $[-3, 3]$. The large, heterogeneity-dependent advantage is therefore a coupling-kernel effect; ownership information accounts at most for the small, implementation-dependent differences in panel (b). Flat-L2 attains the smallest diagnostic in the archived cell because it also changes local normalization, and on the dispatch bases of Section 7.1 the same kernel leaves conditioning essentially unchanged: kernel rankings on synthetic matrices need not transfer.




Fig. 1 Fresh matrix audit. Each exponent range contains 60 generated instances. Panel (a) changes the coupling kernel; panel (b) holds the kernels fixed and removes (Role-Hybrid) or permutes (SA-Mat-permB) the block-scale information. Marks are individual ratios with deterministic horizontal jitter; short horizontal segments are medians. The panels use different vertical scales. Unscaled matrices, whose estimates approach binary64 resolution, enter no ratio; all paired scaled diagnostics are below 6.5×10^6 .

7.3 Completed-run gains are not failure-robust

On the 14 SG-Ter-Mer instances that every policy completed, Flat-GM used less solver work than Base, with paired geometric-mean ratios 0.698 at 1% gap and 0.682 at 0.1% (bootstrap intervals $[0.486, 0.944]$ and $[0.470, 0.925]$): about 30–32% less work on that pool. The fifteenth instance, instance077, completes under Base and times out under every scaled policy at both gaps. Under PAR10 this single failure dominates: Base averages 116 and 134 seconds, while every scaled policy averages between 1,216 and 1,218 seconds (Online Resource 1, Table S4). Scaling therefore produced a sizeable completed-run efficiency gain in this class, but the gain was not robust to failure-sensitive aggregation. On the six instances of the class that are free of the scenario defect the conditional gain is larger and, since instance077 is not among them, the reversal disappears (Online Resource 1, Table S3). The Simple class has no uniform ordering: Flat-GM lowers mean PAR10 from 11.9 to 2.1 seconds at 1% gap and raises it from 52.0 to 73.0 seconds at 0.1%. Role-Hybrid stays within one second of Flat-GM in all four cells.

Replication under a second seed and a second solver. The reconstruction is one solver at one seed. The same 30 instances, four policies, two gaps and 1,800-second limit were rerun changing one factor at a time (Online Resource 1, Table S5). Under Gurobi with seed 7 the SG-Ter-Mer effort ratio of Flat-GM is 0.778 at a 1% gap and 0.777 at 0.1%, against 0.698 and 0.682 at seed 42: the conditional gain survives the change of seed. Under CPLEX at the original seed, in its own deterministic effort metric, the same ratio is 1.078 and 1.077: it does not survive the change of solver. On the Simple class the three conditions agree in direction, Flat-GM below one at the loose gap and between 0.84 and 1.06 at the tight one, Flat-L2 above one in all six cells. The instance that




Fig. 2 Verified optima by coefficient multiplier. Each cell counts the configurations whose returned point is a verified optimum, out of the 48 in that cell (three families, eight seeds, two presolve settings). The fixed margin in (7) brings the continuous feasibility boundary increasingly close to an integer activity. Hatched cells marked – are abstentions, where the policy declines to export. The detached bottom row is GCD-floor, the later integer-aware reformulation control, not a scaling policy.

reverses the ordering, instance077, behaves differently in each condition: under Gurobi at seed 7 every policy times out, Base included, while under CPLEX the run that reaches the limit is Base at the 1% gap and Role-Hybrid at 0.1%. The completed-run gain of this class is therefore a property of a representation together with a pipeline, and the failure that reversed the PAR10 ordering is not a stable property of the instance either.

The two endpoints answer different questions: completed-run effort describes how a representation affects the runs that finish, and PAR10 whether it changes which runs finish. For the reconstructed campaign itself, with one seed and 15 fixed instances per class, the result is a measured conditional gain, not a population estimate of acceleration; the two sensitivity reruns reinforce that caution rather than providing population-level replication. The same runs also show why a stated budget is needed to read residuals: an absolute residual of 4.97 on a zero-RHS equality of caso46e has an activity-normalized value of 1.15×10^{-7} , and without an application budget neither number establishes acceptance; Table 5 supplies one.

7.4 Budgets improve verification; acceptance is pipeline-dependent

These models are small, adversarial and built so that an integer point sits a fixed distance from the continuous boundary, so the counts below compare policies on a deliberately hard family; they are not estimates of how often a solver fails in practice.

The residual budget has a consistent effect within each scaling rule (Table 4). Plain GM centering loses verified optima relative to Base on all three solvers, 26 configurations on HiGHS, 16 on Gurobi and 10 on CPLEX, with no gains. At multipliers 10^6 and 10^9 it shrinks the fixed 0.1 margin of (7) below the solver

Table 4 Paired verification outcomes on identical model, presolve and solver settings (192 per comparison and solver). “A only” counts configurations verified under *A* and not under *B*.

Comparison <i>A</i> vs <i>B</i>	HiGHS		Gurobi		CPLEX	
	<i>A</i> only	<i>B</i> only	<i>A</i> only	<i>B</i> only	<i>A</i> only	<i>B</i> only
Budget-GM vs GM	41	0	12	0	0	0
Budget-dyadic vs Dyadic-GM	19	0	13	0	0	0
Budget-only vs GM	50	0	11	0	0	0
GM vs Base	0	26	0	16	0	10
Budget-only vs Base	24	0	1	6	0	10
Budget-GM vs Base	15	0	1	5	0	10
Budget-dyadic vs Base	2	8	1	4	0	10
GM-tight vs GM	25	0	26	0	26	0
Budget-only vs GM-tight	25	0	4	19	0	26
Budget-only vs Budget-GM	9	0	0	1	0	0

tolerance, and most of its failures (49 of 51 on HiGHS, 36 of 52 on Gurobi, 34 of 52 on CPLEX) return points that satisfy the exported tolerance while violating the original row by more than the budget. Adding the budget to the same rule never lost a verified configuration on any solver. It recovered 41 of them on HiGHS and 12 on Gurobi, and none on CPLEX, where a rule and its budgeted version verify exactly the same configurations; the dyadic pair shows the same pattern, 19, 13 and 0. The two comparisons say different things about the same factors. Plain centering improves the numbers in the exported file and loses verified optima; the same rule with a budget attached loses none. A representation can be numerically better without being operationally safer, unless what safer means is stated in the units of the original model.

Relative to the unscaled model, the outcome depends on the pipeline: presolve, internal scaling, tolerances and finite precision all intervene between the exported row and the accepted point. On HiGHS, Budget-only verifies 191 of 192 configurations and Budget-GM 182, against 167 for Base and 141 for GM; Budget-only gains 24 configurations over Base and loses none. On Gurobi, Budget-only verifies 151, Budget-GM 152 and Base 156. On CPLEX every scaled policy verifies 140 and Base 150: the budget changes nothing there, and the ten configurations lost to scaling are lost with it and without it. On both commercial solvers, tightening the two tolerances to 10^{-9} (GM-tight, 166 configurations on each of the three solvers) is the more effective intervention. The same stored models and row factors thus lead to different acceptance in three pipelines, and the budget helps where this experiment shows the greater sensitivity to the exported coefficients.

The gap between Budget-only and Budget-GM on HiGHS is informative in its own right. At multipliers 10^6 and 10^9 , the only ones at which their outcomes differ, the budget endpoint is binding and Budget-GM selects that endpoint rounded inward, one unit in the last place (ulp) away from Budget-only’s factor. All ten discordant configurations, nine on HiGHS and one on Gurobi, with none on CPLEX, have exactly this one-ulp difference. The nine-configuration gap therefore measures how sensitive solver acceptance is to

the last bit of a factor near a tolerance boundary, not an effect of coefficient centering. The same sensitivity appears at a larger scale: Budget-dyadic, which uses the nearest admissible power of two and exports every product exactly, verifies 161 configurations on HiGHS, 153 on Gurobi and 140 on CPLEX. Exact representation of the scaled row does not by itself secure acceptance.

Two controls delimit the result. GCD division with RHS flooring verifies all 192 configurations on each solver: it removes the deliberately narrow integer margin, whereas a budget controls the original-unit meaning of an exported residual, so the budget result is a comparison among scalings and not an advantage over integer-aware presolve. Budget-round, which reserves the worst-case rounding term of (6), verifies all 48 configurations it accepts per solver and abstains on the other 144, including many that other policies solve correctly. Its abstentions certify that the requested uniform contract cannot be asserted; they are not solved instances, and Figure 2 counts them separately.

The same contract on the application. The 0.01 of the stress family is a test parameter, but the dispatch rows carry interpretable units, so a budget can be stated in them. For illustration, suppose the modeller declares 0.1 MW for the rows that account for generation and cap it by demand, 0.1 MW for the piecewise-linear production envelopes, and 1,000 m³ — a thousandth of a cubic hectometre — for the reservoir water balances. Table 5 applies Proposition 5.1 to the 4,320 rows of caso46e that carry those meanings, with $\epsilon = 10^{-6}$ and the coefficient window $[10^{-9}, 10^9]$. The admissible interval is wide, the geometric-mean factor lies inside it on every row, and the residual the tolerance allows in original units stays five to nine orders of magnitude below the declared budget. On this model the contract leaves the geometric-mean factor unchanged on every row: it restricts which factors are admissible, as it must, but the geometric-mean choice considered here already lies well inside that restriction. The budgets are declared rather than measured, so the margin is also the sensitivity: scaling all six by a common multiplier, all three kernels stay admissible on all 4,320 rows down to 10^{-4} and the two the comparison turns on down to 10^{-5} , which places the result far from the chosen numbers (Online Resource 1, Table S7).

Its content is in the reading of a residual. One family is written pre-scaled: the Baygorria generation definition carries a coefficient of 10^7 on the generation variable, so the same 0.1 MW becomes a budget of 10^6 in the units of that row. It is the family on which the accepted Flat-L2 run of Section 7.3 leaves its largest residual, 4.97, which against the row's own budget is a generation error of 5×10^{-7} MW. The residual as written and the quantity it stands for differ by the seven orders of magnitude of that coefficient, and only a budget carried with the row distinguishes them.

Table 5 A declared residual contract on the dispatch model, applied to the 4,320 rows of caso46e whose units are interpretable. Budgets are stated by the modeller, not measured or supplied by the operator. Admissible factors are the interval of Proposition 5.1 for $\epsilon = 10^{-6}$ and coefficient limits $[10^{-9}, 10^9]$; the last column is ϵ divided by the geometric-mean factor, the original-unit residual that the tolerance allows, in the units of that group’s budget. The last group is written pre-scaled, with a coefficient of 10^7 on the generation variable, so 0.1 MW there is 10^6 in the units of the row as written. The admissibility check covers the three kernels the four policies use — the unit factor, the geometric-mean factor shown here and $1/\|a_r\|_2$ — and every one of them is admissible on every row. The archived admissibility check records the corresponding counts for each kernel.

Constraint group	Budget δ	Admissible factors	GM factor	Residual at ϵ
Generation identities, demand cap	0.1 MW	$[10^{-5}, 10^9]$	1	10^{-6}
Generation definition, Bonete/Palmar	0.1 MW	$[10^{-5}, 10^9]$	1	10^{-6}
Production envelope tangents	0.1 MW	$[10^{-5}, 10^9]$	2.14	5.0×10^{-7}
Water balance, Bonete/Palmar	10^3 m^3	$[10^{-9}, 2.8 \times 10^5]$	0.0167	6.0×10^{-5}
Water balance, Baygorria	10^3 m^3	$[10^{-9}, 2.8 \times 10^5]$	2.8×10^{-4}	3.6×10^{-3}
Generation definition, Baygorria	10^6 row units	$[7.0 \times 10^{-8}, 100]$	2.6×10^{-3}	3.8×10^{-4}

8 Discussion

8.1 What we learned

Row scaling can substantially change numerical conditioning. With the basis held fixed, the representation alone moves the condition number by six to eight orders of magnitude in the median, and the kernel decides whether it does.

The mechanism is simpler than the metadata explanation. Under exact GM local normalization, the extrema-based post-local block statistic carries no block-scale information (Proposition 4.1), so the block-relative rule becomes Euclidean normalization of the coupling rows. Separately, on a locally whitened triangular structure the optimal common factor is $(1 + \rho^2)^{-1/2}$ (Proposition 4.3); the two coincide when the structure has a single coupling row.

Matching kernels is essential for attribution. The original comparison changed the coupling kernel together with the information used, and the fixed-basis audit shows how much a kernel alone can change: identical coverage, representations orders of magnitude apart. A claim that generator information improves preprocessing needs a kernel-matched control.

Residual budgets are a meaningful invariant quantity. A residual budget carried consistently with its row defines a contract that is invariant under consistent re-emission — the number changes with the row, the meaning does not — and determines the admissible factors exactly. In the stress test, measured against the corresponding unbudgeted rule, it repaired the verification losses of coefficient centering on two of the three solvers and cost nothing on the third; measured against the unscaled model the outcome still depends on the pipeline. On the dispatch model the same contract is easy to satisfy — all three kernels the four policies use are admissible on each of its 4,320 interpretable rows, with

orders of magnitude to spare — and its work is in the reading of a residual, which without a stated unit says nothing. GCD division with RHS flooring verified every stress configuration, but it answers a different question: it is an integer-aware reformulation that exploits the activity lattice of integer rows, whereas a residual budget is a contract in original units that also applies to continuous rows and to constraints admitting no such reduction.

Solver acceptance and effort are properties of the pipeline. The same representation reduced completed-run work and caused a timeout, the same factors were accepted differently by the three pipelines, and a one-ulp change in a factor changed outcomes in one of them. Pipeline, not solver: presolve, internal scaling, tolerances and finite precision all sit between the exported row and the point that comes back, and the experiments here cannot separate their contributions.

8.2 Scope and limitations

The operational evidence comes from one application: two instance classes and 15 fixed instances per class, from the reconstructed campaign of Section 6.3. The two sensitivity reruns, under a second seed and a second solver, are enough to show that at least one finding is configuration-dependent rather than a stable property of the representation, and far too small to describe the population of configurations. The stress instances are small and deliberately adversarial; their budget of 0.01 is a numerical test parameter, and the dispatch budgets of Table 5 are stated by us to show what the contract implies, not supplied or validated by the operator. Solver outcomes depend on the solver, its version and its settings. The collapse result concerns the extrema-based block statistic after GM normalization; other block statistics, other kernels and physical scale information are outside it, and on the operational models the coupling rows receive the same factors under every GM-local policy, so ownership is tested only on synthetic matrices. Finally, row scaling cannot repair within-row heterogeneity beyond the limit in (3) or the integer-rounding amplification in (6).

9 Conclusion

External row scaling genuinely matters numerically. Holding one basis fixed, it changes the conditioning of generated dispatch models by six to eight orders of magnitude in the median while preserving the MILP in exact arithmetic, and the kernel determines whether this happens.

Correct attribution changes the explanation of the gain. The advantage originally attributed to block metadata is a coupling-kernel effect: exact GM local normalization erases the extrema-based post-local block statistic, and the closed-form coupling factor of the idealized triangular model, which reduces to

Euclidean normalization in the single-row case, provides a spectral explanation for it that uses no ownership information.

Residual budgets provide a constructive way to encode information that keeps its meaning in original application units. They come with exact feasibility and incompatibility conditions, a centered factor, invariance under re-emission, and explicit rounding and binary64 export conditions, and in the stress experiment adding the residual budget to the same scaling rule never reduced verification against that same unbudgeted rule, on any of the three solvers tested.

What a solver does with a scaled representation still has to be measured, and its output still has to be verified. The practical lesson is therefore not to avoid external scaling, but to make it accountable: attribute its gains with matched controls, state its accuracy contract in the model's own units, and verify solver output against that contract.

Data and software availability

The accompanying repository contains the C++ generator, frozen historical aggregates, the reconstructed operational logs and analysis, the fixed-basis audit with its exports, bases and independent verification, the exact-enumeration experiment, the row-factor implementation, tests, and the scripts that generate the tables and vector figures. The current research sources are distinguished from the historical draft, whose positive block-attribution claim is superseded. All newly reported data are stored under `results-revision/research-audit/`. No commercial solver is required to regenerate the figures, rerun the fixed-basis audit, verify the saved binary solutions, or run the HiGHS stress arm. Rerunning the Gurobi and CPLEX arms requires the corresponding licenses.

The repository is at <https://github.com/MathiasRodriguezCastro/structure-aware-row-scaling>. The state reviewed here is the fixed release `v1.0-mpc-submission`, which carries three archives: the manuscript sources, the source and data archive that is Online Resource 2, and the bulk inputs of the fixed-basis audit, 5.18 GB of exported LP text and bases in a 337 MB compressed file. Only the rebuild of all 1,530 bases from that text needs the last one; every number reported here follows from the summaries in Online Resource 2.

Supplementary information

Online Resource 1 provides the triangular-model proof, the fixed-basis audit protocol and complete ranges, operational reconstruction details, frozen experiment protocols, arithmetic implementation checks, and reproduction instructions. The accompanying source and data archive is Online Resource 2.

References

1. Achterberg, T., Bixby, R.E., Gu, Z., Rothberg, E., Weninger, D.: Presolve reductions in mixed integer programming. *INFORMS Journal on Computing* **32**(2), 473–506 (2020). <https://doi.org/10.1287/ijoc.2018.0857>
2. Berthold, T., Hendel, G.: Learning to scale mixed-integer programs. *Proceedings of the AAAI Conference on Artificial Intelligence* **35**(5), 3661–3668 (2021). <https://doi.org/10.1609/aaai.v35i5.16482>
3. Cheung, K.K.H., Gleixner, A., Steffy, D.E.: Verifying integer programming results. In: F. Eisenbrand, J. Koenemann (eds.) *Integer Programming and Combinatorial Optimization, Lecture Notes in Computer Science*, vol. 10328, pp. 148–160. Springer (2017). https://doi.org/10.1007/978-3-319-59250-3_13
4. Curtis, A.R., Reid, J.K.: On the automatic scaling of matrices for Gaussian elimination. *IMA Journal of Applied Mathematics* **10**(1), 118–124 (1972). <https://doi.org/10.1093/imamat/10.1.118>
5. Demmel, J.: Nearly optimal Block-Jacobi preconditioning. *SIAM Journal on Matrix Analysis and Applications* **44**(1), 408–413 (2023). <https://doi.org/10.1137/22M1504901>
6. Elble, J.M., Sahinidis, N.V.: Scaling linear optimization problems prior to application of the simplex method. *Computational Optimization and Applications* **52**(2), 345–371 (2012). <https://doi.org/10.1007/s10589-011-9420-4>
7. GAMS Development Corporation: Special features for mathematical programs: Model scaling—the scale option. GAMS User’s Guide, online documentation (2026). URL https://www.gams.com/latest/docs/UG_LanguageFeatures.html. Accessed September 8, 2026
8. Gurobi Optimization, LLC: Tolerances and user-scaling. Gurobi Optimizer Reference Manual, online documentation (2026). URL https://docs.gurobi.com/projects/optimizer/en/current/concepts/numericguide/tolerances_scaling.html. Accessed September 8, 2026; page updated June 17, 2026
9. Higham, N.J., Tisseur, F.: A block algorithm for matrix 1-norm estimation, with an application to 1-norm pseudospectra. *SIAM Journal on Matrix Analysis and Applications* **21**(4), 1185–1201 (2000). <https://doi.org/10.1137/S0895479899356080>
10. Hoen, A., Gleixner, A.: Analyzing the numerical correctness of branch-and-bound decisions for mixed-integer programming. In: G. Tack (ed.) *Integration of Constraint Programming, Artificial Intelligence, and Operations Research, Lecture Notes in Computer Science*, vol. 15763, pp. 35–50. Springer (2025). https://doi.org/10.1007/978-3-031-95976-9_3
11. Horn, S.: How do I scale my model to improve its numerical properties? Gurobi Help Center (2025). URL <https://support.gurobi.com/hc/en-us/articles/26284047605521-How-do-I-scale-my-model-to-improve-its-numerical-properties>. Updated December 22, 2025; accessed September 8, 2026

12. IDAES Project: CustomScalerBase class. IDAES v2.12.0 documentation (2026). URL https://idaes-pse.readthedocs.io/en/stable/reference_guides/scaling/custom_scaler.html. Accessed September 8, 2026
13. IDAES Project: Scaling methods. IDAES v2.12.0 documentation (2026). URL https://idaes-pse.readthedocs.io/en/stable/reference_guides/core/util/scaling.html. Accessed September 8, 2026
14. IEEE: IEEE Standard for Floating-Point Arithmetic. IEEE (2019). <https://doi.org/10.1109/IEEESTD.2019.8766229>. IEEE Std 754-2019
15. Jalving, J., Shin, S., Zavala, V.M.: A graph-based modeling abstraction for optimization: Concepts and implementation in Plasmo.jl. *Mathematical Programming Computation* **14**, 699–747 (2022). <https://doi.org/10.1007/s12532-022-00223-3>
16. Klotz, E.: Identification, assessment, and correction of ill-conditioning and numerical instability in linear and integer programs. In: *Bridging Data and Decisions, INFORMS TutORials in Operations Research*, pp. 54–108. INFORMS (2014). <https://doi.org/10.1287/educ.2014.0130>
17. Pritchett, R., Howell, K., Grebow, D.: Low-thrust transfer design based on collocation techniques: Applications in the restricted three-body problem. In: *AAS/AIAA Astrodynamics Specialist Conference*. Stevenson, Washington (2017). URL https://engineering.purdue.edu/people/kathleen.howell.1/Publications/Conferences/2017_AAS_PriHowGre.pdf. AAS Paper 17-626
18. Pyomo Development Team: ScaleModel: Model scaling transformation. Pyomo 6.9.1 documentation (2026). URL <https://pyomo.readthedocs.io/en/6.9.1/api/pyomo.core.plugins.transform.scaling.ScaleModel.html>. Accessed September 8, 2026
19. Qu, Z., Gao, W., Hinder, O., Ye, Y., Zhou, Z.: Optimal diagonal pre-conditioning. *Operations Research* **73**(3), 1479–1495 (2025). <https://doi.org/10.1287/opre.2022.0592>

Acknowledgements

The experiments reported in this paper were carried out using ClusterUY, the national supercomputing facility of Uruguay (<https://cluster.uy>).

Statements and Declarations

Funding. The author received no external funding, grant or material support for this work.

Competing interests. The author declares no financial or non-financial competing interests. No solver vendor or software project discussed here had any role in the design, execution, analysis or reporting of the study.

Author contributions. Mathias Rodríguez Castro is the sole author and is responsible for the conception of the study, the implementation, the experiments, the analysis and the manuscript. The model generator used as the source of instances originates in his hydrothermal-dispatch thesis work.

Ethics approval, consent and human or animal participants. Not applicable: the study involves no human participants, animals or personal data.

Use of AI assistance. AI assistants (OpenAI Codex and Anthropic Claude) were used as coding and writing tools during the preparation of this work: for implementation and verification code, for derivations and their review, for analysis of the literature, and for drafting and editorial revision. Reported numerical results come from the stored computational runs, and figures were generated programmatically from those data or the stated equations. The author independently verified the mathematical arguments, the experimental design, the computational results and the scientific conclusions, and takes full responsibility for the manuscript.